



# CS 524 A

Introduction to Cloud Computing

Lecture 13:

Identity Management and Security in the Cloud, Part II

# AGENDA

- Authentication
  - Authentication protocol principles
  - Authentication with Kerberos
  - More on authentication with PKI
  - IPsec and TLS
- Authorization (OAuth as an example)
- Identity Management in the Cloud
  - Requirements
  - OpenStack IdM (*Keystone*)
  - Use of TPM (secure boot, attestation, etc.)
  - Use of HSM

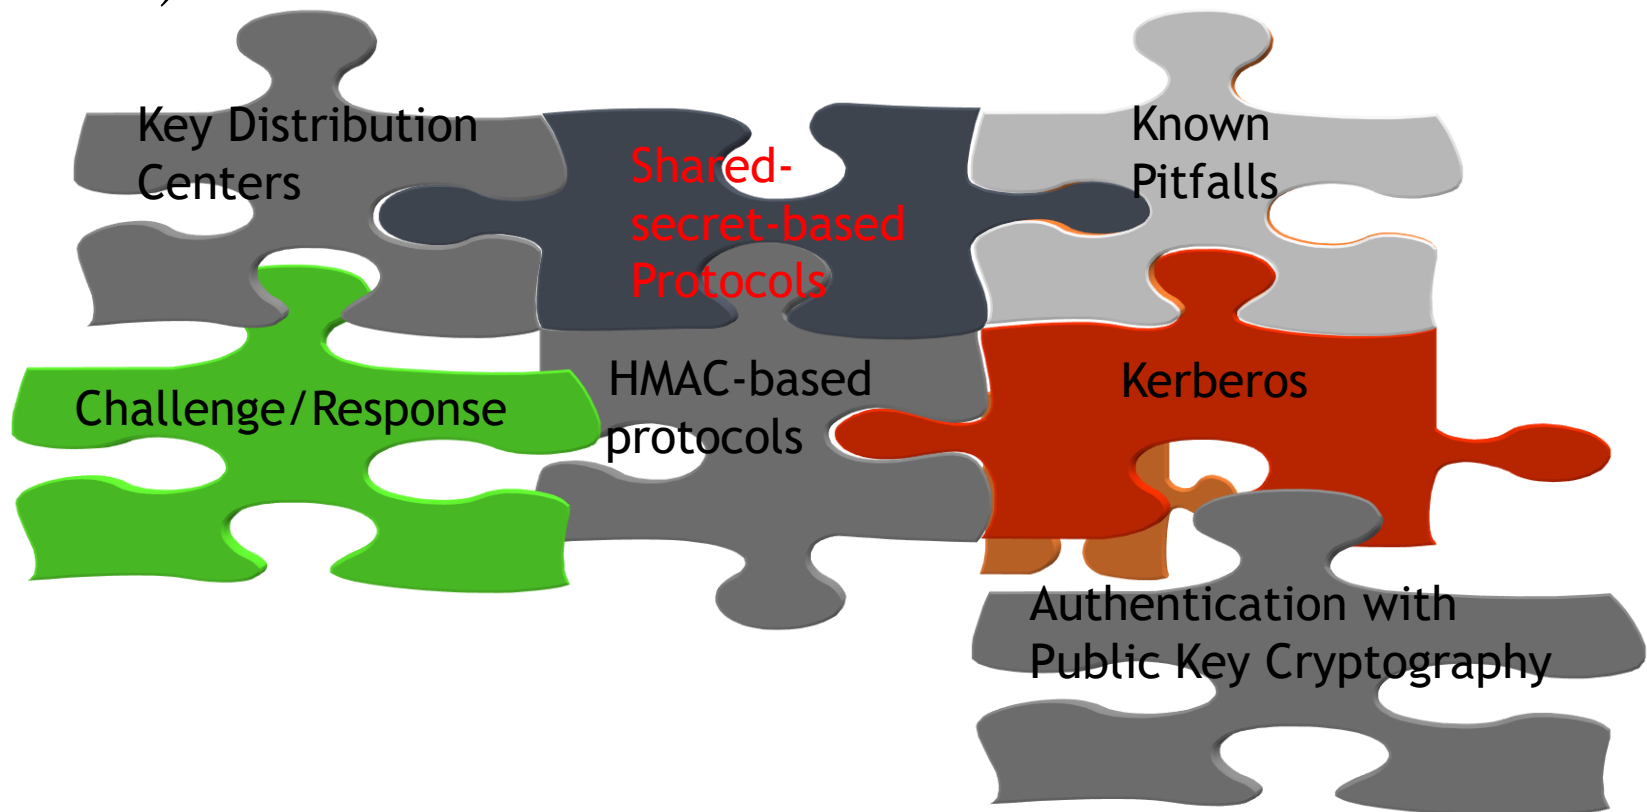


# BACK TO BASIC SECURITY PROPERTIES

- Confidentiality
  - Keeping information secret from unintended users
  - Achieved through encryption*
- Authentication
  - Confirming the identity of the presenter of the information
  - Achieved by authentication protocols*
- Authorization
  - Determining whether a user may have access to a resource
  - Achieved through access control lists (ACLs), implementing policies, and also by token-base authorization protocols (OAuth, OpenStack Keystone tokens, etc.)*
- Integrity
  - Ensuring that a message received was the one that was actually sent
  - Achieved by cryptographic means called signatures*
- Non-repudiation
  - Ensuring that a party that has signed a contract cannot later deny having signed it
  - Achieved by either third-party attestation or asymmetric cryptographic means*

# AUTHENTICATION PROTOCOLS

Needed for the establishment of sessions (VoIP conversations [streams and signaling], TCP sessions, etc.)



# INTRODUCTION TO KEY PLAYERS

Alice



Bob



Trudy the Intruder



# THE GENERAL MODEL

- Alice starts by sending a message to Bob
- An exchange follows
- Trudy may intercept, modify, or replay any message

# CHALLENGE-RESPONSE PROTOCOL (FIRST ATTEMPT)

Alice

Bob

A Identifier

$R_B$  Challenge: A *nonce*--a large random number, not to be repeated

$K_{AB}(R_B)$  Response, encrypted with the shared--or private--key

$R_A$

$K_{AB}(R_A)$

# CHALLENGE-RESPONSE PROTOCOL (CAN WE DO THIS FASTER?)

Alice

Bob

$A, R_A$

$R_B, K_{AB}(R_A)$

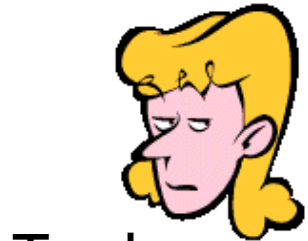
$K_{AB}(R_B)$

An improvement:  
3 instead of 5 messages!

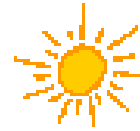
No!



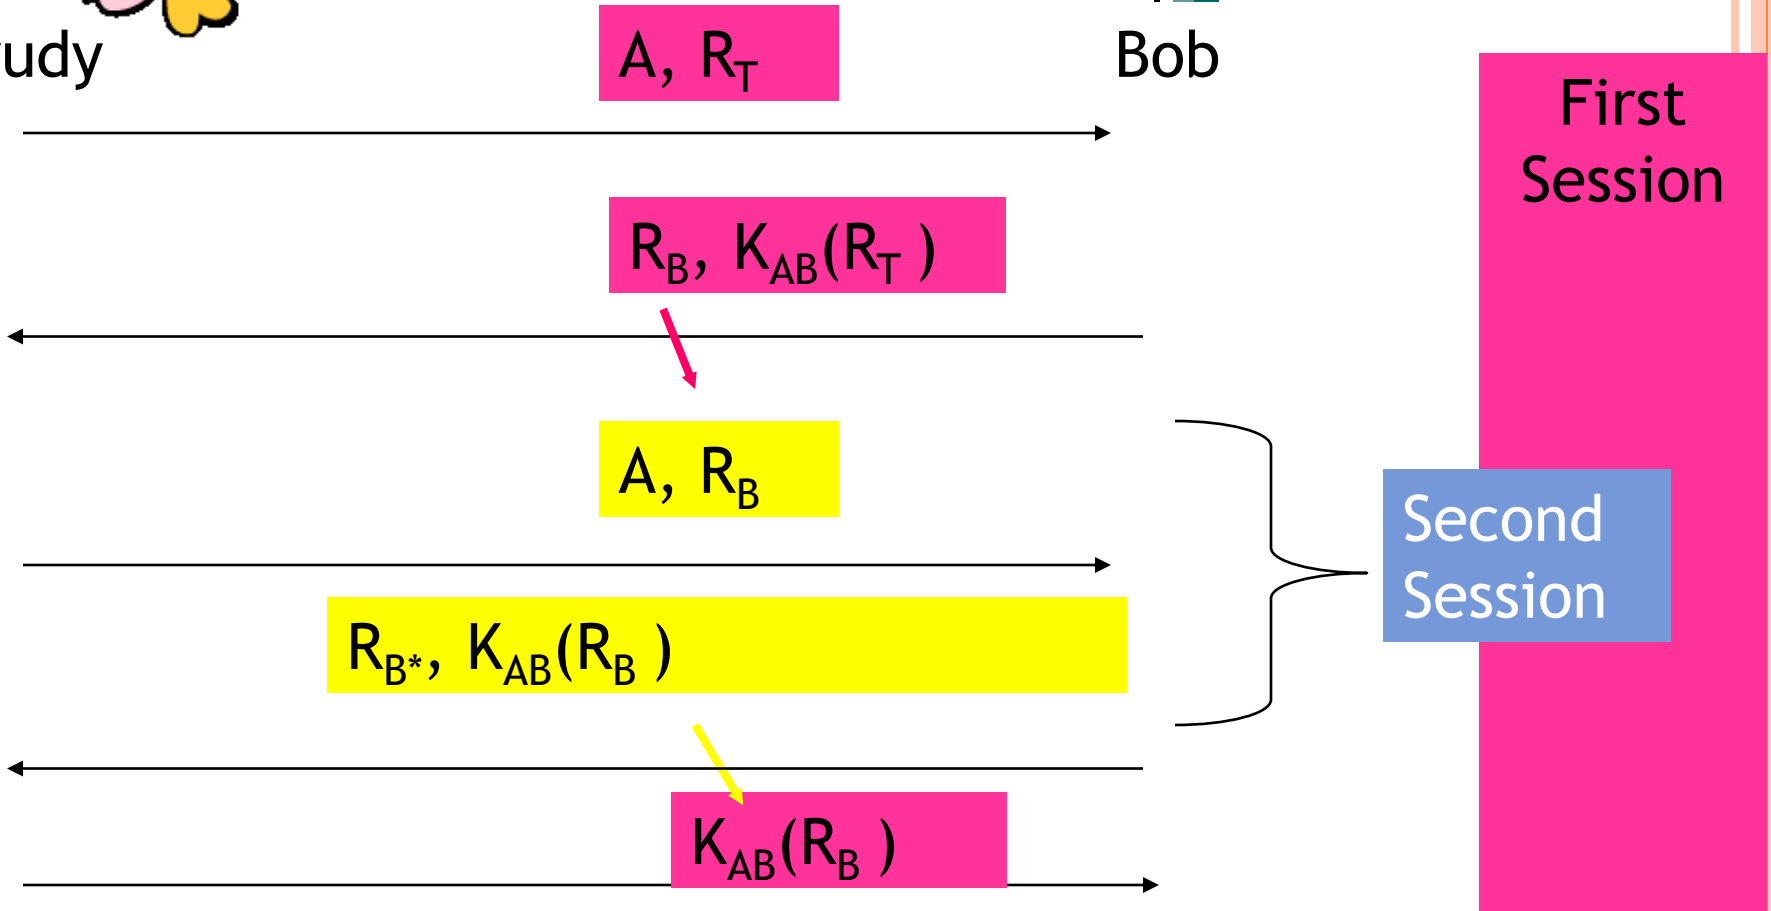
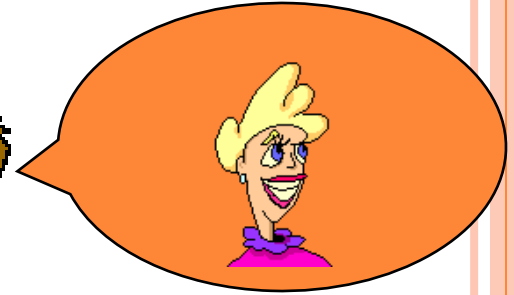
# THE REFLECTION ATTACK



Trudy



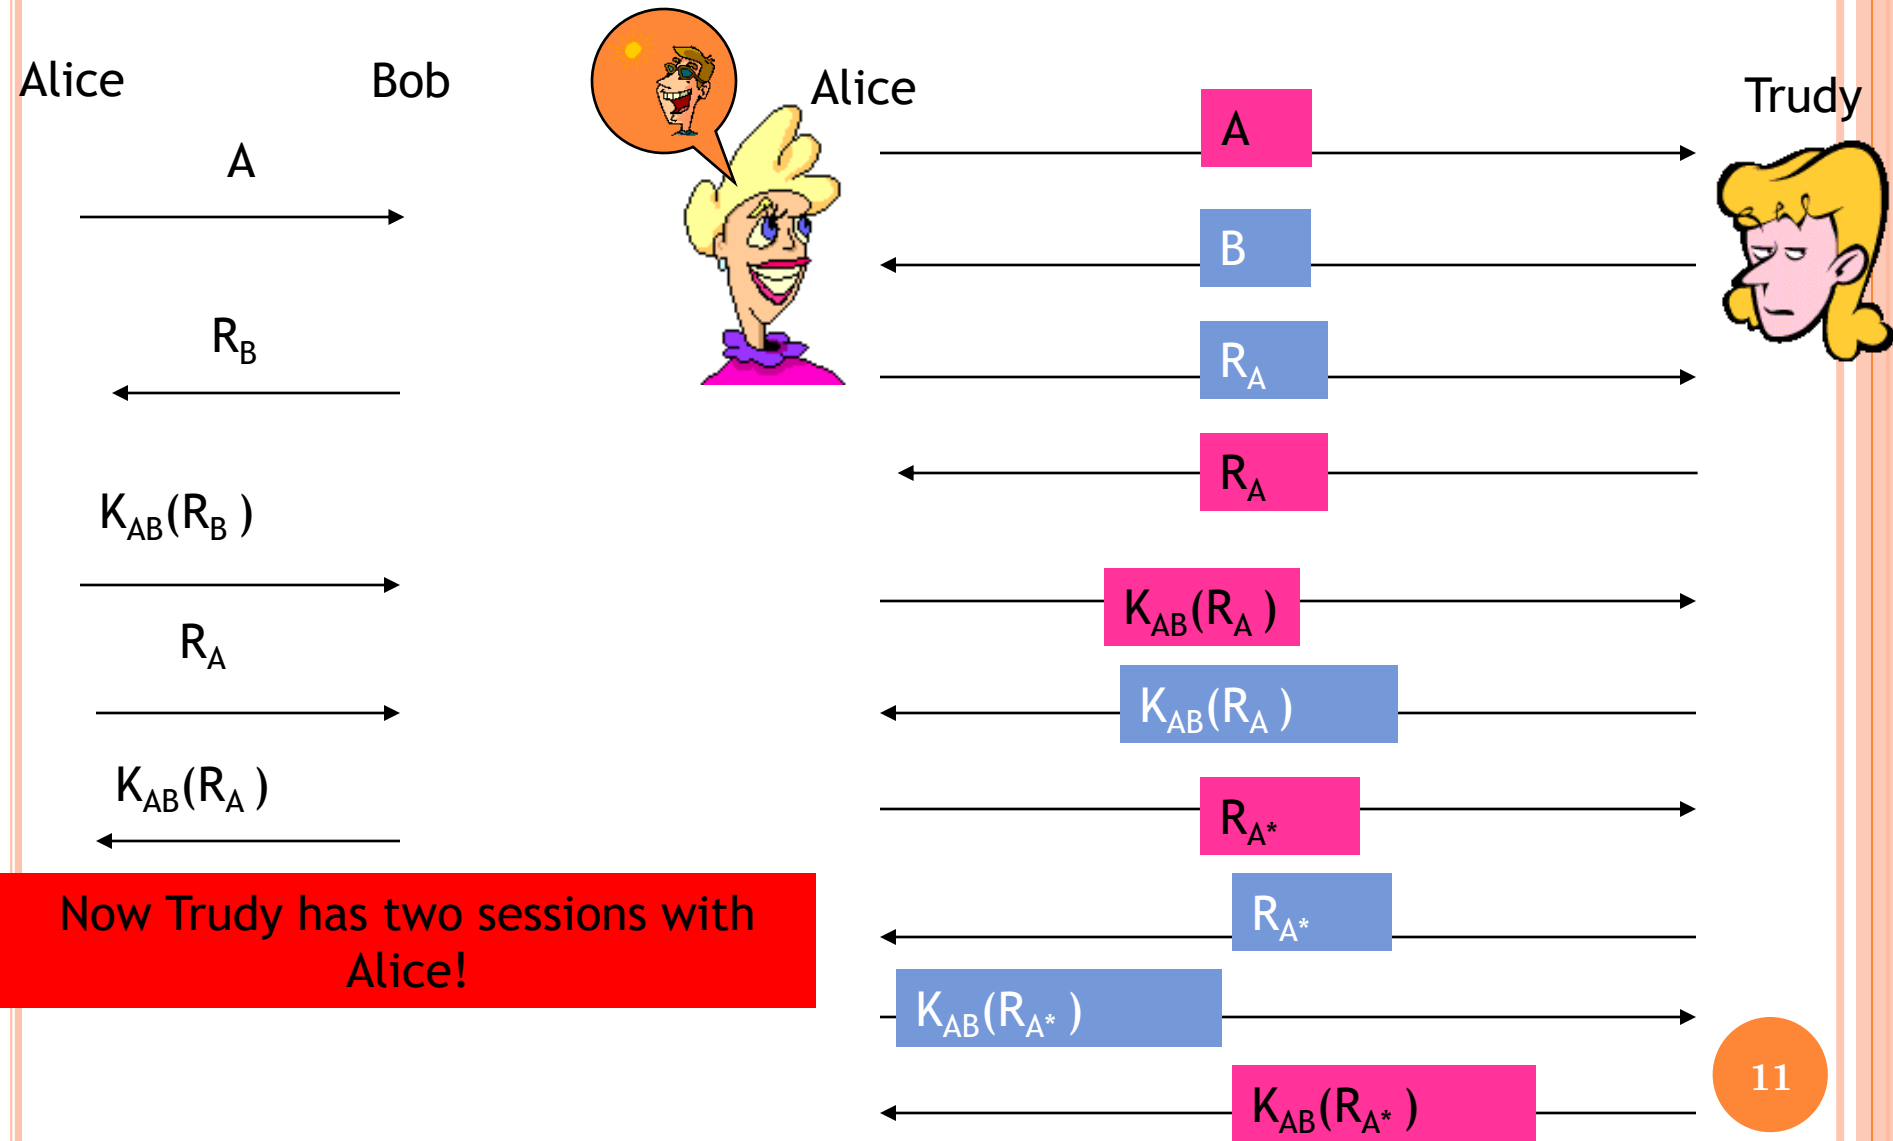
Bob



# GENERAL (TEXTBOOK) RULES

- The initiator has to prove its identity before the responder
- The initiator and responder must use different keys for proof (a need for two shared keys)
- Initiator and responder must draw challenges from different sets (e.g., odd/even)
- It must be impossible to use authentication information obtained in one session in a different one

# BUT WAS THE FIRST ATTEMPT REALLY FAULTLESS?



# A FEW CONCLUSIONS

- The authentication protocols are very hard to design correctly
- There is a method of designing protocols of this kind that are *provably* correct: R. Bird et al., *Systematic Design of a Family of Attack-Resistant Authentication Protocols*, IEEE Journal on Selected Areas in Communications, vol. 11, pp. 679-693, June 1993

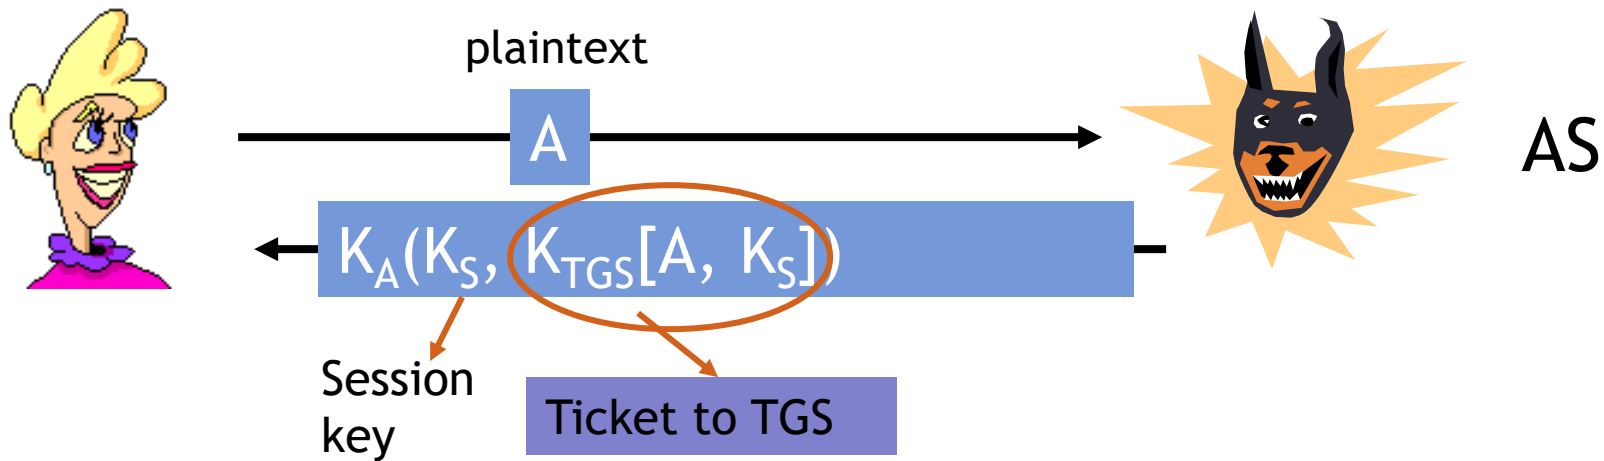
# AUTHENTICATION WITH KERBEROS

- Kerberos was originally developed by MIT based on a variant of *Needham-Shroeder* protocol and then standardized in the IETF
  - RFC 4120 specifies Kerberos V5
- Kerberos assumes that all clocks are synchronized
- Kerberos modifies the KDC model

# THE KERBEROS MODEL: THREE SERVERS

- Authentication Server (AS)
  - Authenticates users during the *login* session
  - Shares a secret (password) with every user
- Ticket-Granting Server (TGS)
  - Issues proof-of-identity tickets, which convince other servers that the owners of the tickets are who they claim to be
- The Application Server (AP)
  - Does the real work (performs services such as banking transactions, telephone calls, etc.)

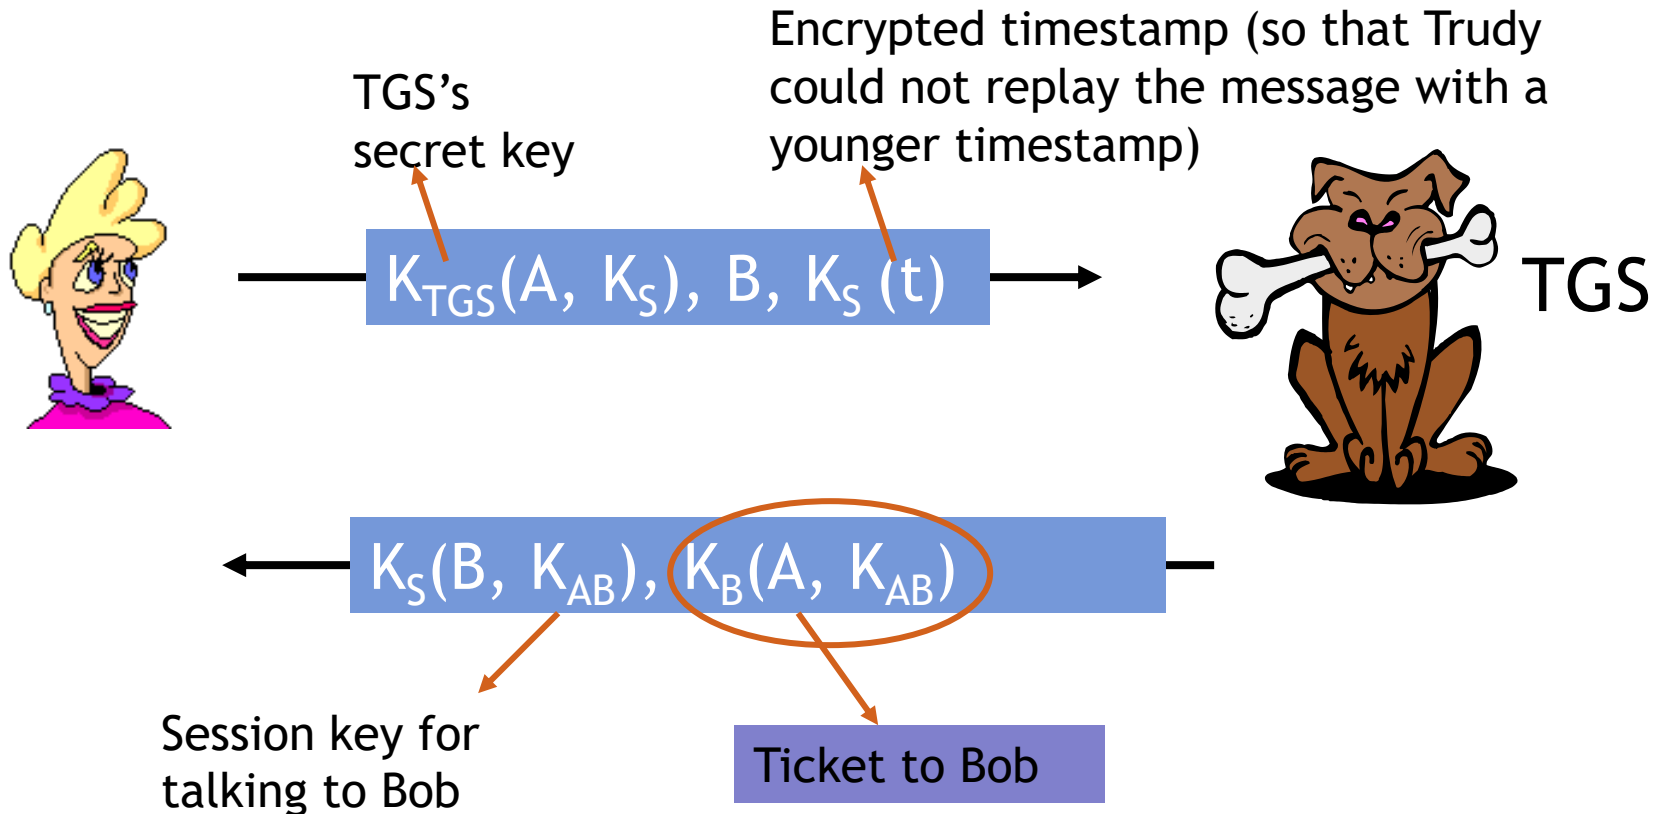
# OPERATION OF KERBEROS AS



At this point,

- 1) Alice is prompted for a password by the client, and this password is used for generating  $K_A$ , so that she obtains the session key and the ticket for TGS
- 2) The client *forgets* the password
- 3) Alice says she wants to use Bob's services

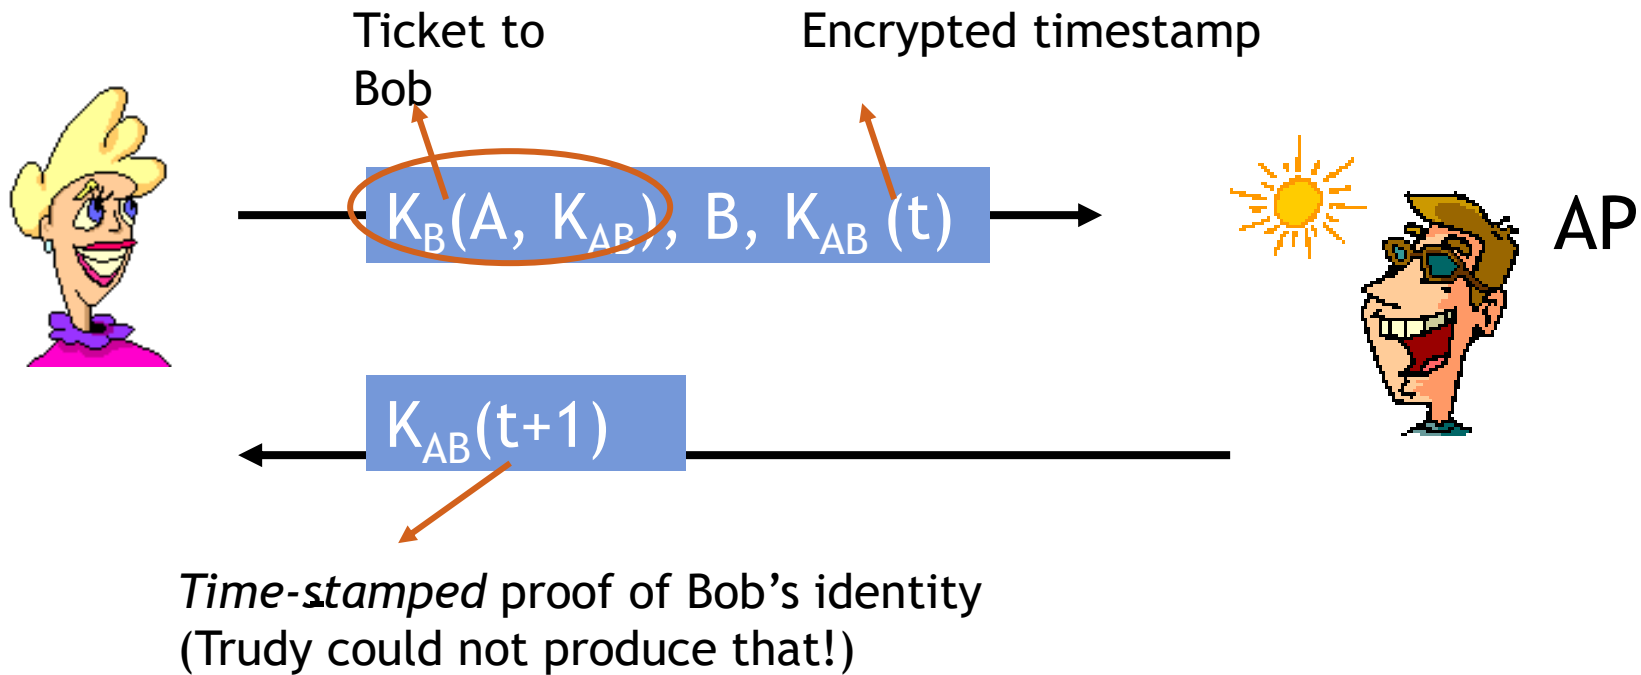
# OPERATION OF KERBEROS TGS



Now Alice can start talking to the real-work server—Bob

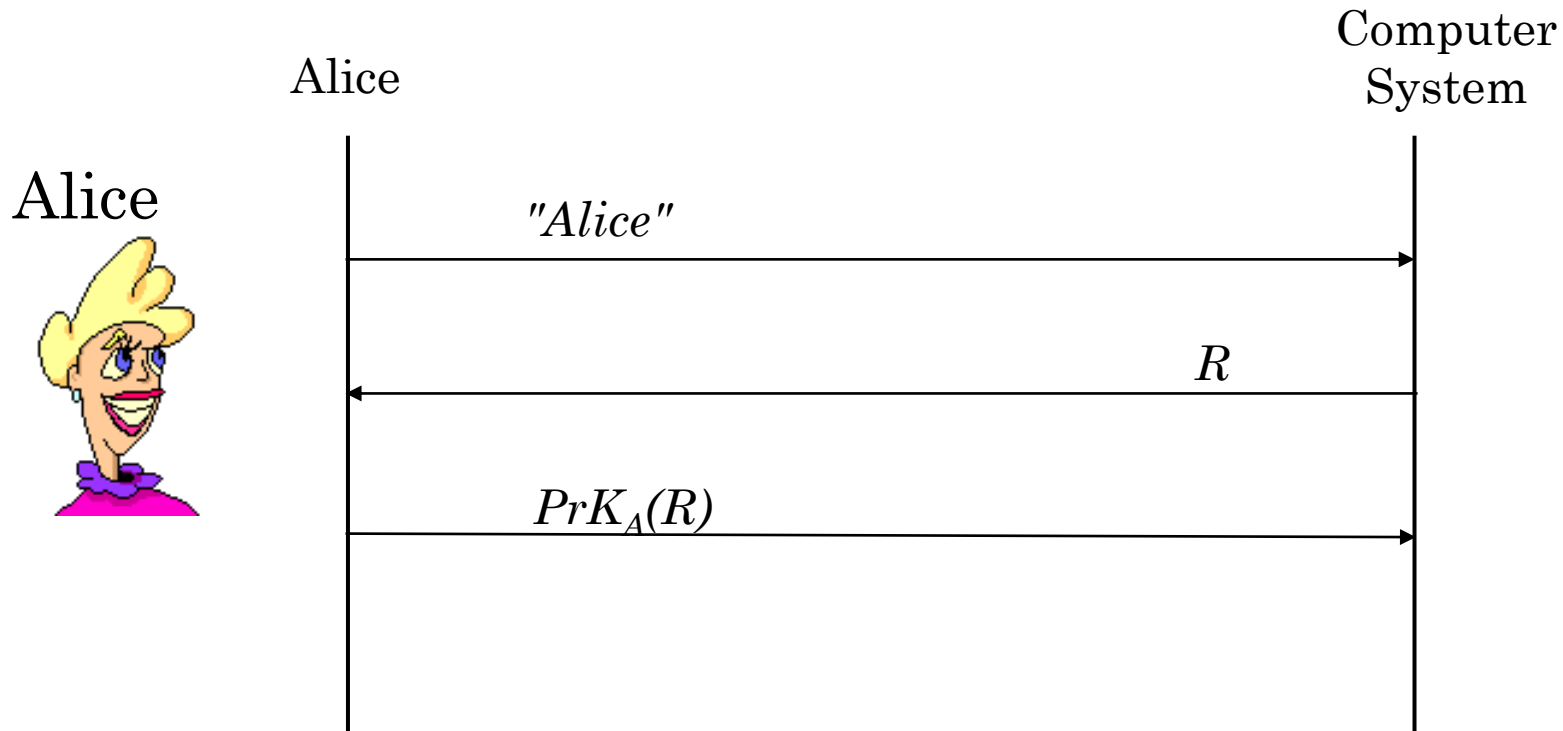


# OPERATION OF KERBEROS APPLICATION SERVER

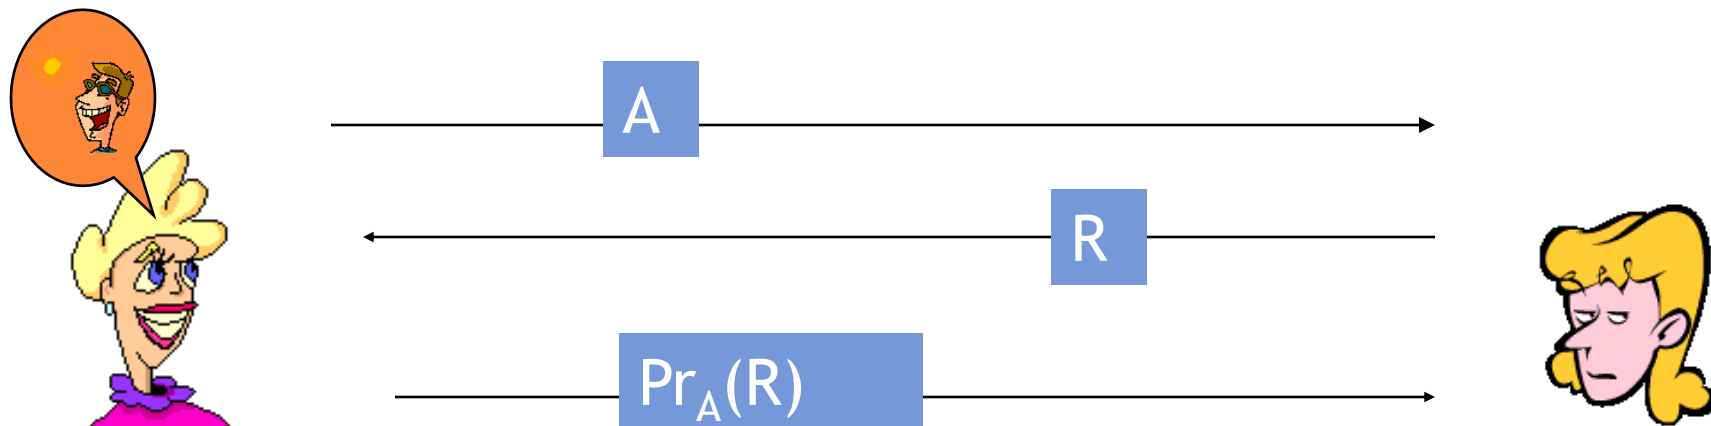


**Now Alice can work with Bob, but if she needs to change to another application server, she just restarts with the request to TGS (no passwords are ever transmitted)**

# PUBLIC-KEY-BASED AUTHENTICATION



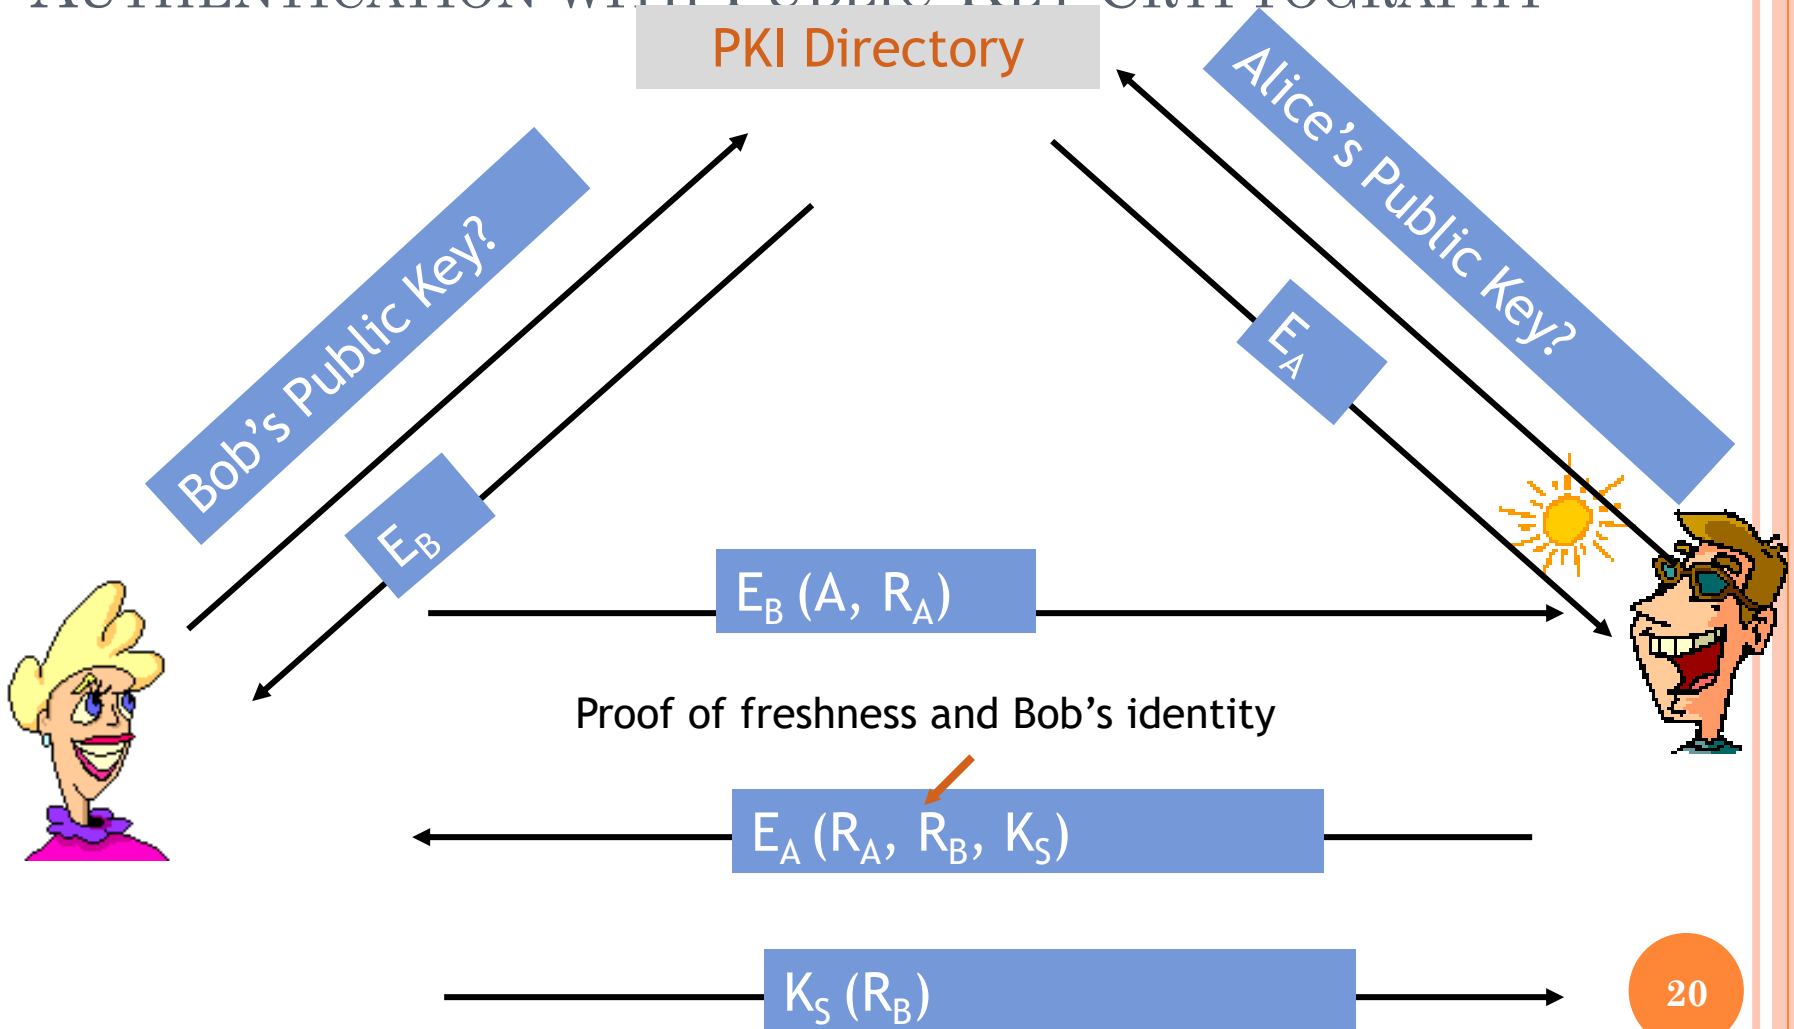
# AUTHENTICATION WITH PUBLIC-KEY CRYPTOGRAPHY: A NAÏVE “SOLUTION”



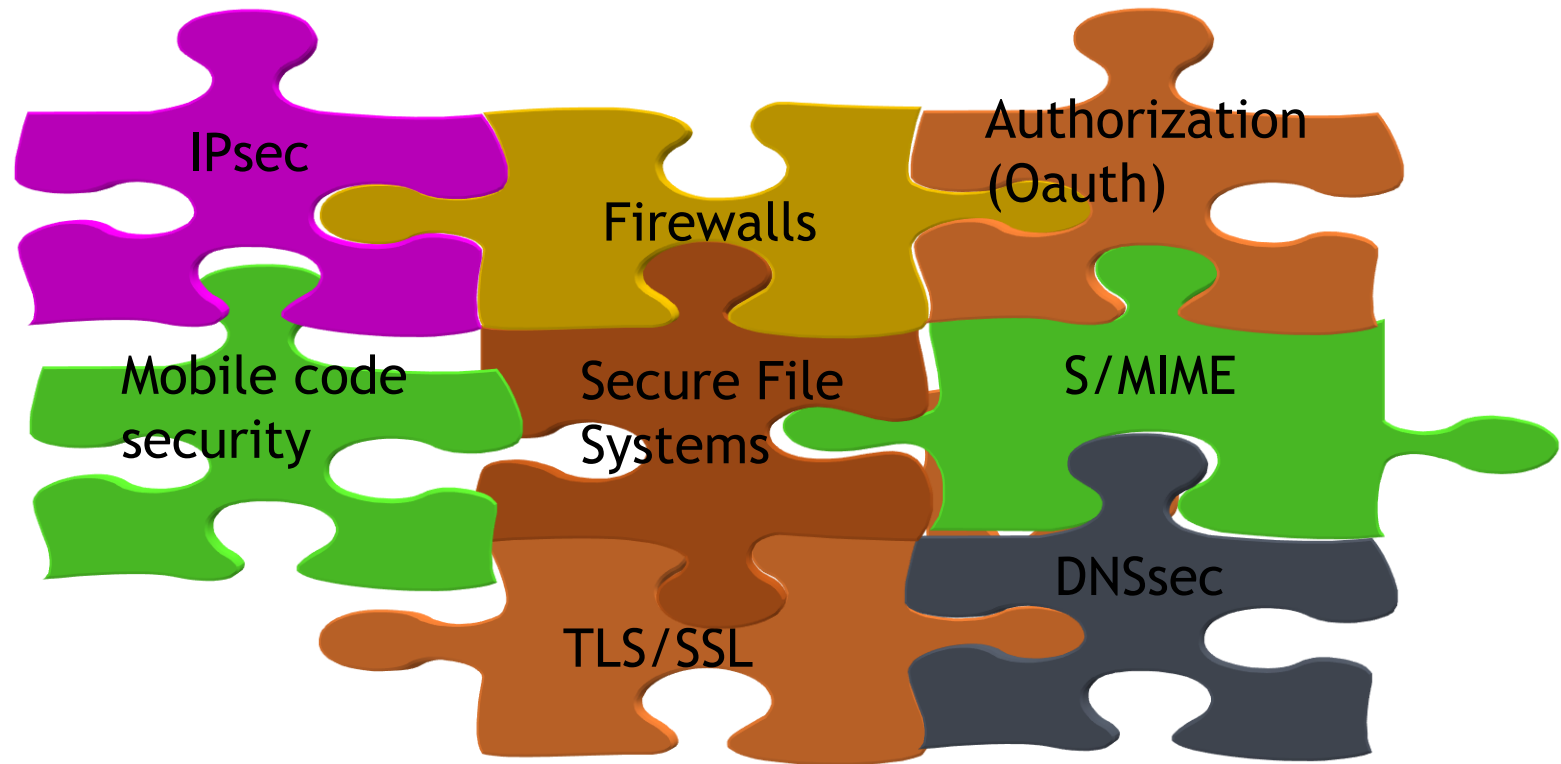
... $R$  = “I, undersigned Alice, owe Trudy \$1,000,000”

# AUTHENTICATION WITH PUBLIC-KEY CRYPTOGRAPHY

PKI Directory



# COMMON SECURITY SOLUTIONS USED IN CLOUD



We will focus on *IPsec* and *TLS*, which have become indispensable for authentication and *Oauth*, which is considered for authorization

# NETWORK SECURITY IN THE PROTOCOL STACK

Application Layer

Application-specific protocols

Presentation Layer

Encrypt the whole session

Session Layer

Transport Layer

Network Layer: firewalls help with limited success

Link Layer: nothing needs to be done, if it is really point-to-point (unlike *radio*); otherwise, use link encryption

Physical Layer: Prevent wiretapping by enclosing transmission lines in sealed tubes containing argon at high pressure monitored by an alarm

# TWO VIEWS IN THE INTERNET CAMP

Security must be end-to-end, and for this reason alone must be implemented in the Application Layer (which will make plaintext unavailable to operating systems)

Problem: Then all applications must be re-written and... how many people really understand security to rewrite them?

Security must be implemented in the Network Layer without users ever approaching it!

Problem: Even though this view has prevailed, a truly network-layer implementation proved to be impossible, and Internet principles had to be violated.

# IP SECURITY PROTOCOL (*IPSEC*)

*IPsec* is a framework for multiple

- Services
  - confidentiality, integrity, protection from replay—among the major ones
- Algorithms
  - to make it algorithm-independent (and there is a *Null algorithm*)
- Granularities
  - from a single TCP connection to an aggregate

*IPsec* is... connection-oriented!



# SECURITY ASSOCIATION (SA)

- SA is a *simplex* connection identified by Security Parameters Index (SPI) carried by all packets
- SA is needed because
  - a key must be used for some period of time—the duration of the connection
  - the set up time is amortized among many packets

# ESTABLISHING AN SA

This involves

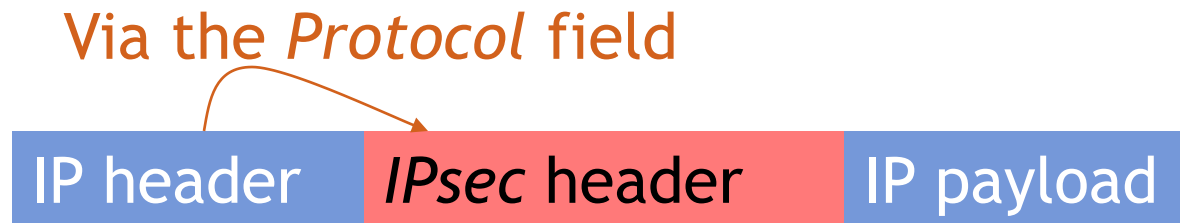
- Authenticating both ends
- Establishing the key
- Agreeing on cryptographic algorithms
- Initializing the sequence number (which will run through the life of the association)
- Establishing SPI

## TWO PARTS OF *IPSEC*

- The *Internet Security Association and Key Management Protocol (ISAKMP)* deals with establishing symmetric keys
  - The main protocol is called *Internet Key Exchange (IKE)*. It has problems, and it is being replaced by *IKE2*
- The other part deals with the headers defined for the two modes of *IPsec* operation
  - Transport mode and
  - Tunnel mode

# TRANSPORT- AND TUNNEL MODES

- Transport mode



- Tunnel mode



Useful for

- terminating at other than end-user locations (e.g., firewalls)
- aggregation to prevent *traffic analysis*

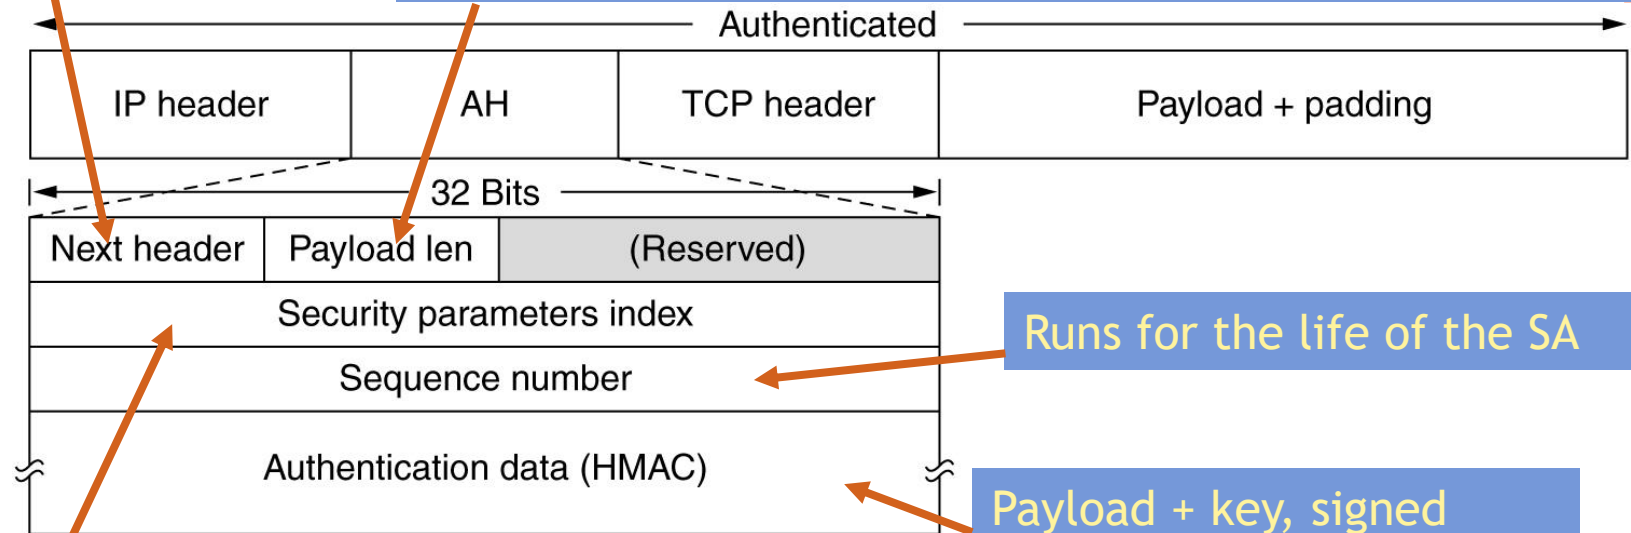
## TWO (HISTORICAL) HEADERS

- The *Authentication Header (AH)* deals only with integrity checking but *not* confidentiality; hashed message authentication code (HMAC) covers only immutable IP fields (not TTL)
- The *Encapsulating Security Payload (ESP)* supports both HMAC integrity and full confidentiality. In a way, it makes AH superfluous
- Both, of course, provide authentication

# AUTHENTICATION HEADER (AH) (IPv4 TRANSPORT MODE)

Stores the value that *IP Protocol* field had

Number of 32-bit words in AH minus 2



Runs for the life of the SA

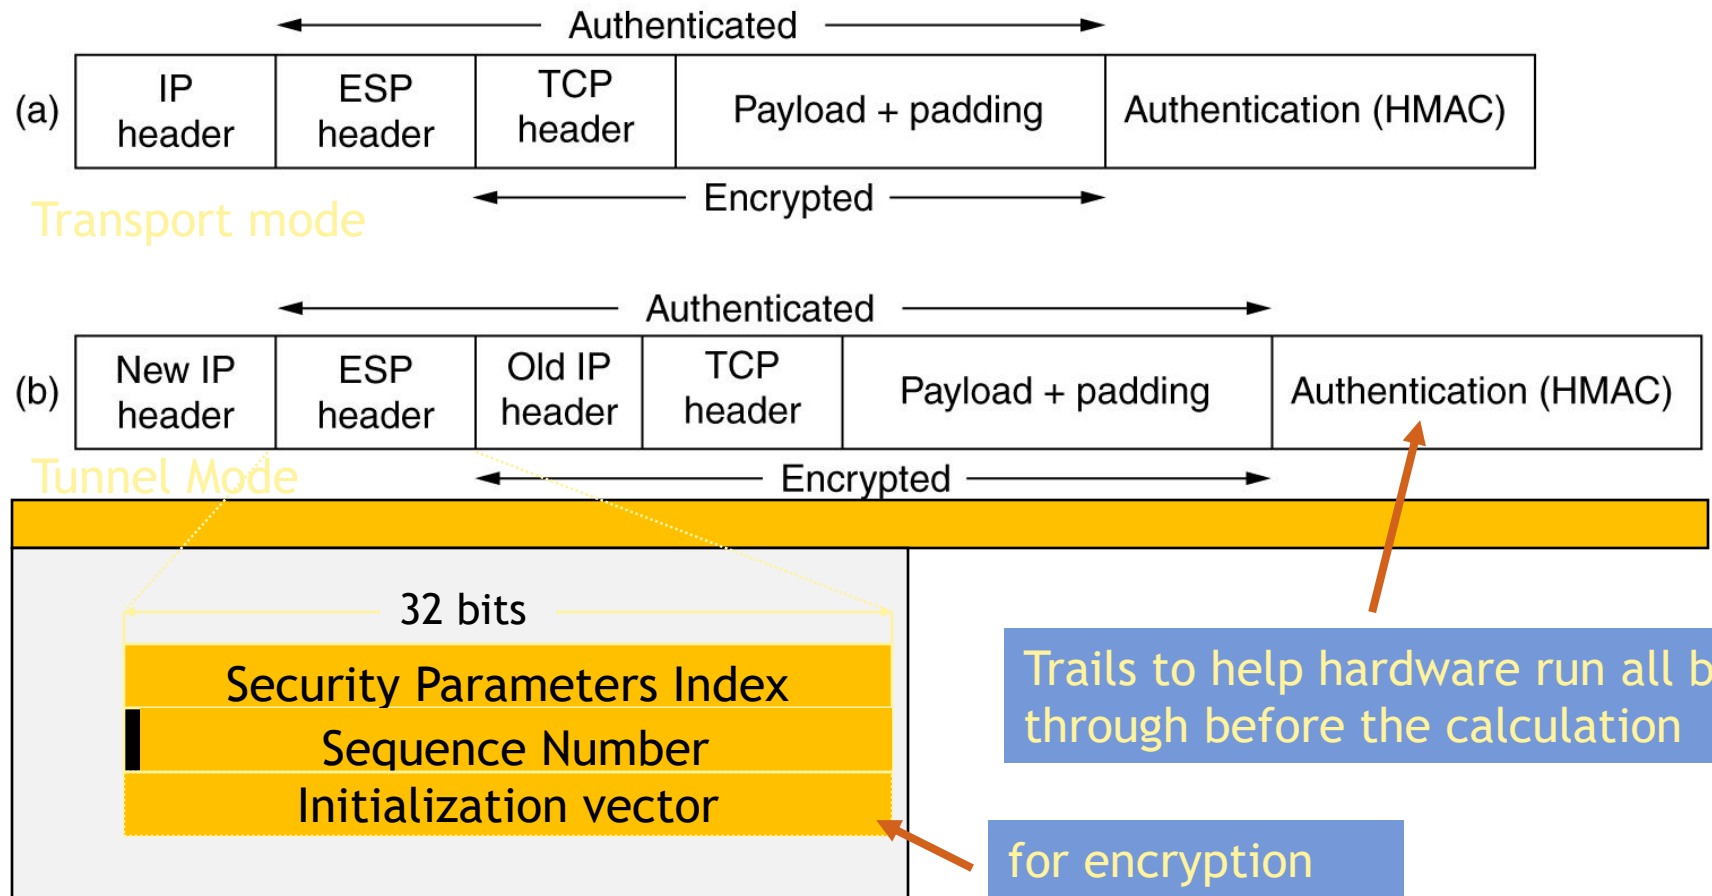
Payload + key, signed

The “virtual circuit number” associated with the shared key

(After A. Tanenbaum)

30

# ENCAPSULATING SECURITY PAYLOAD (ESP) HEADER



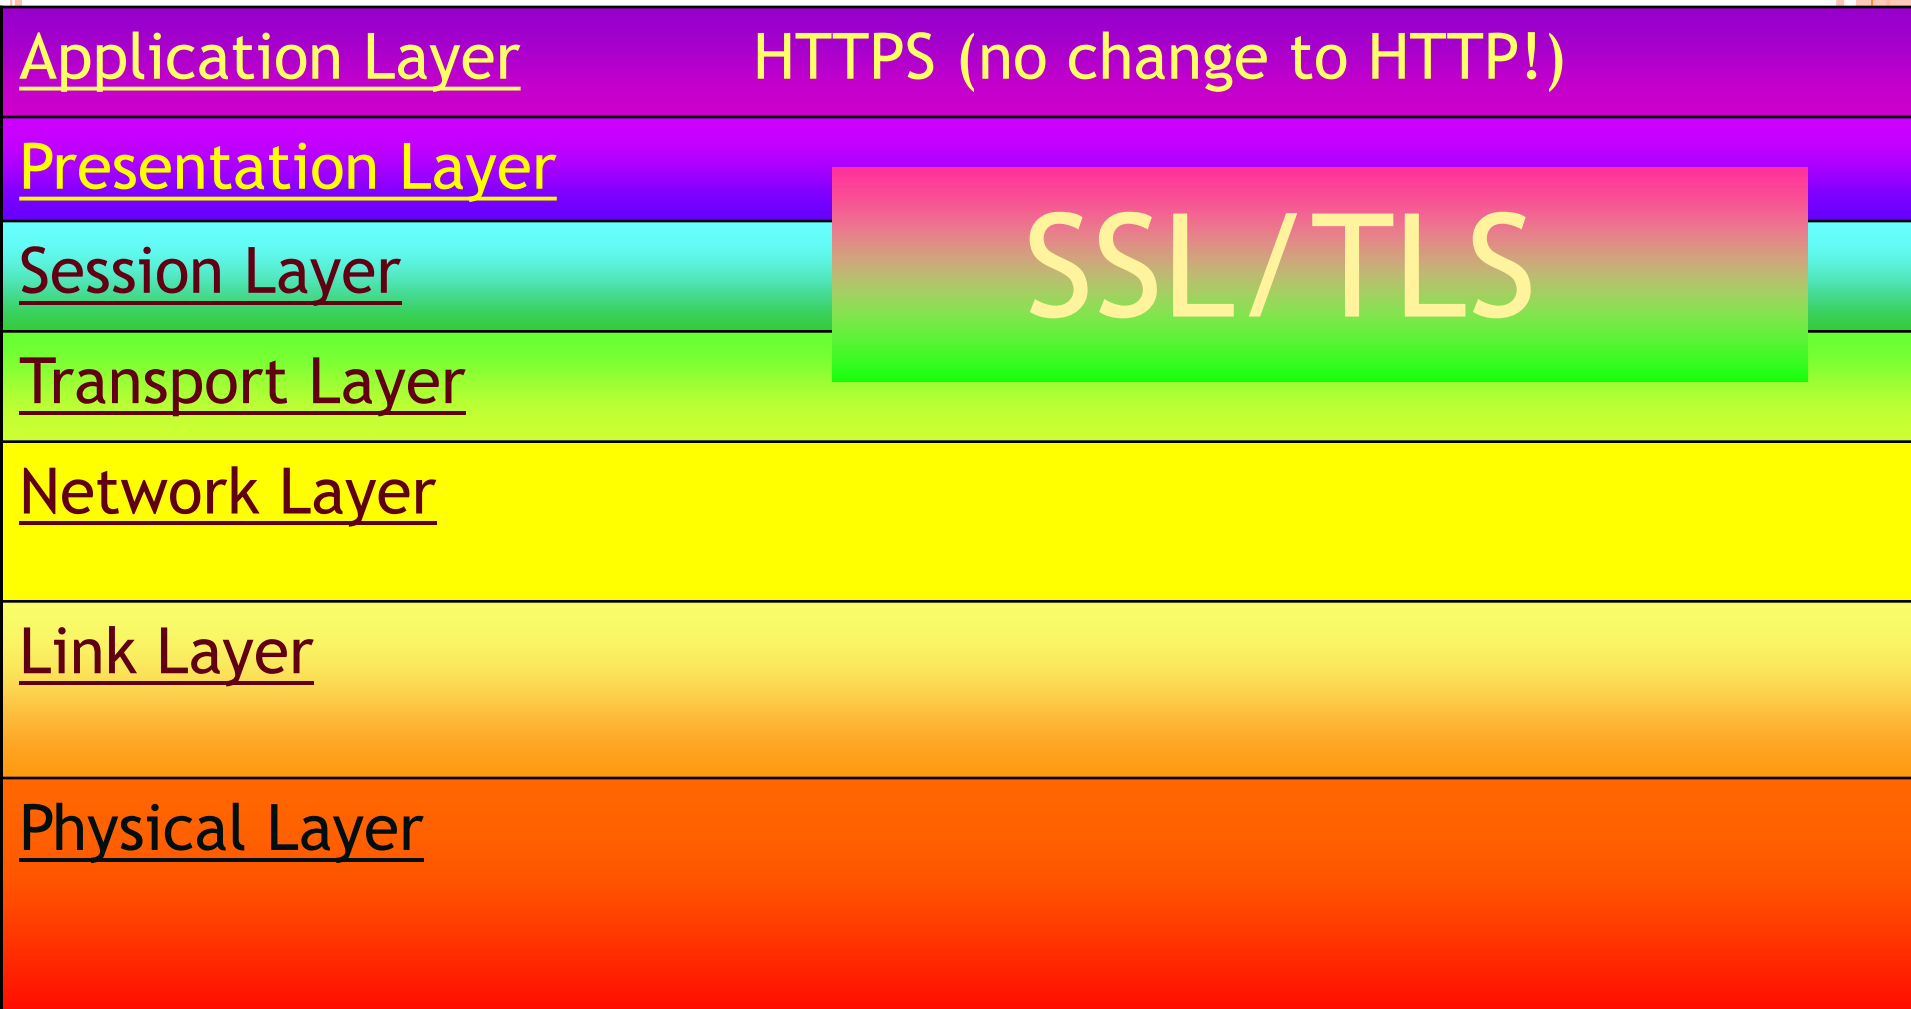
(After A. Tanenbaum)

# TRANSPORT LAYER SECURITY (TLS) BACKGROUND

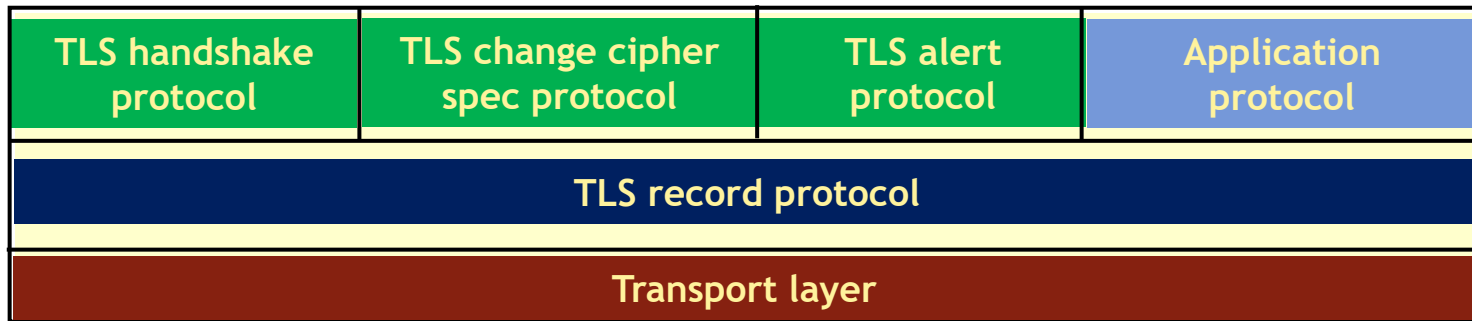
- TLS has been designed to
  - protect a *client-server* session
  - run over TCP (or, actually, any reliable transport protocol)
  - be used by communicating process at a *user-level* (rather than *system-level* so as to avoid operating system changes)
- It has been developed by the IETF based on the *Netscape* Secure Socket Layer (SSL) v3
- The latest version is TLS v1.2 (RFC 5246), which supersedes the previous two versions v1.0 (RFC 2246) and v1.1 (RFC 4346)



# POSITION OF THE SSL/TLS IN THE OSI REFERENCE ARCHITECTURE



# TLS ARCHITECTURE



## Application protocol

**TLS *handshake* protocol:** Signaling for session initiation, including authentication and negotiation of session parameters (e.g., random numbers and the cipher suite)

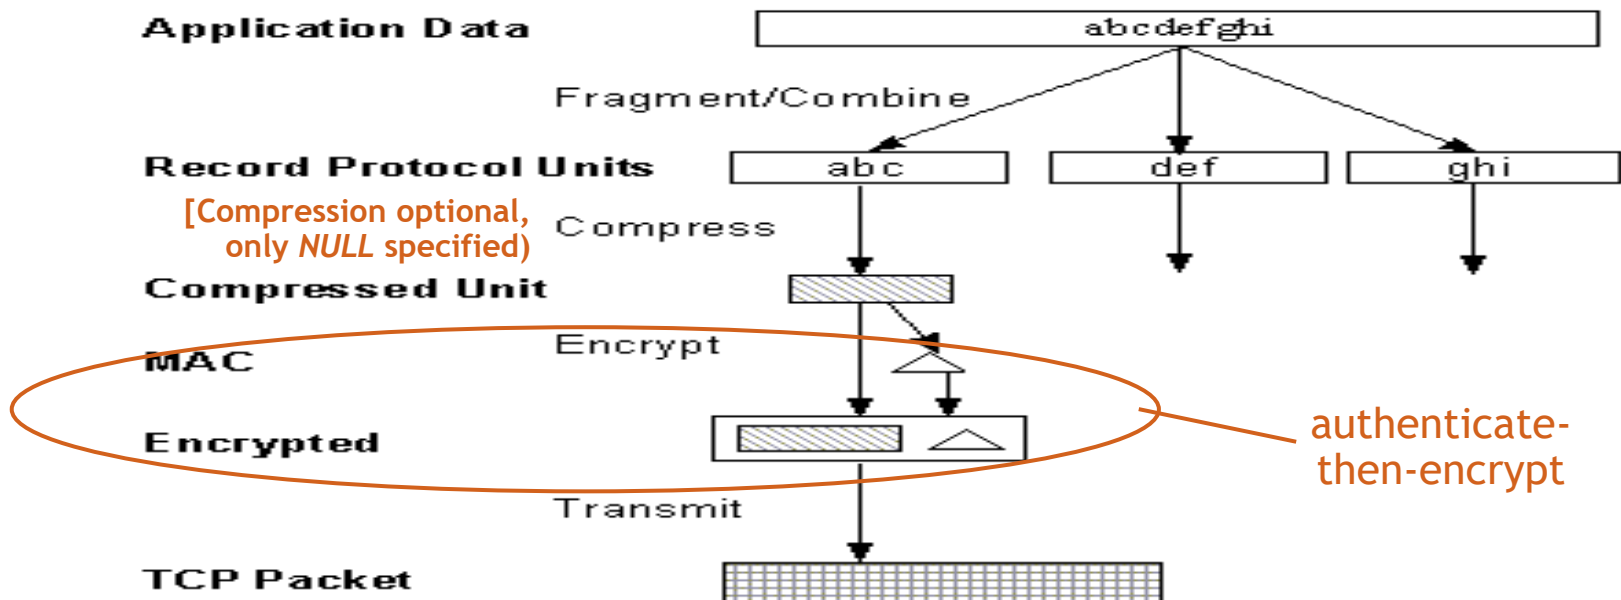
**TLS *change cipher spec* protocol:** Signaling to use the newly negotiated cipher suite and session keys for protecting ensuing records

**TLS *alert* protocol:** Signaling errors during a session

**TLS *record* protocol:** Processing application data for transmission, including digest computation and encryption

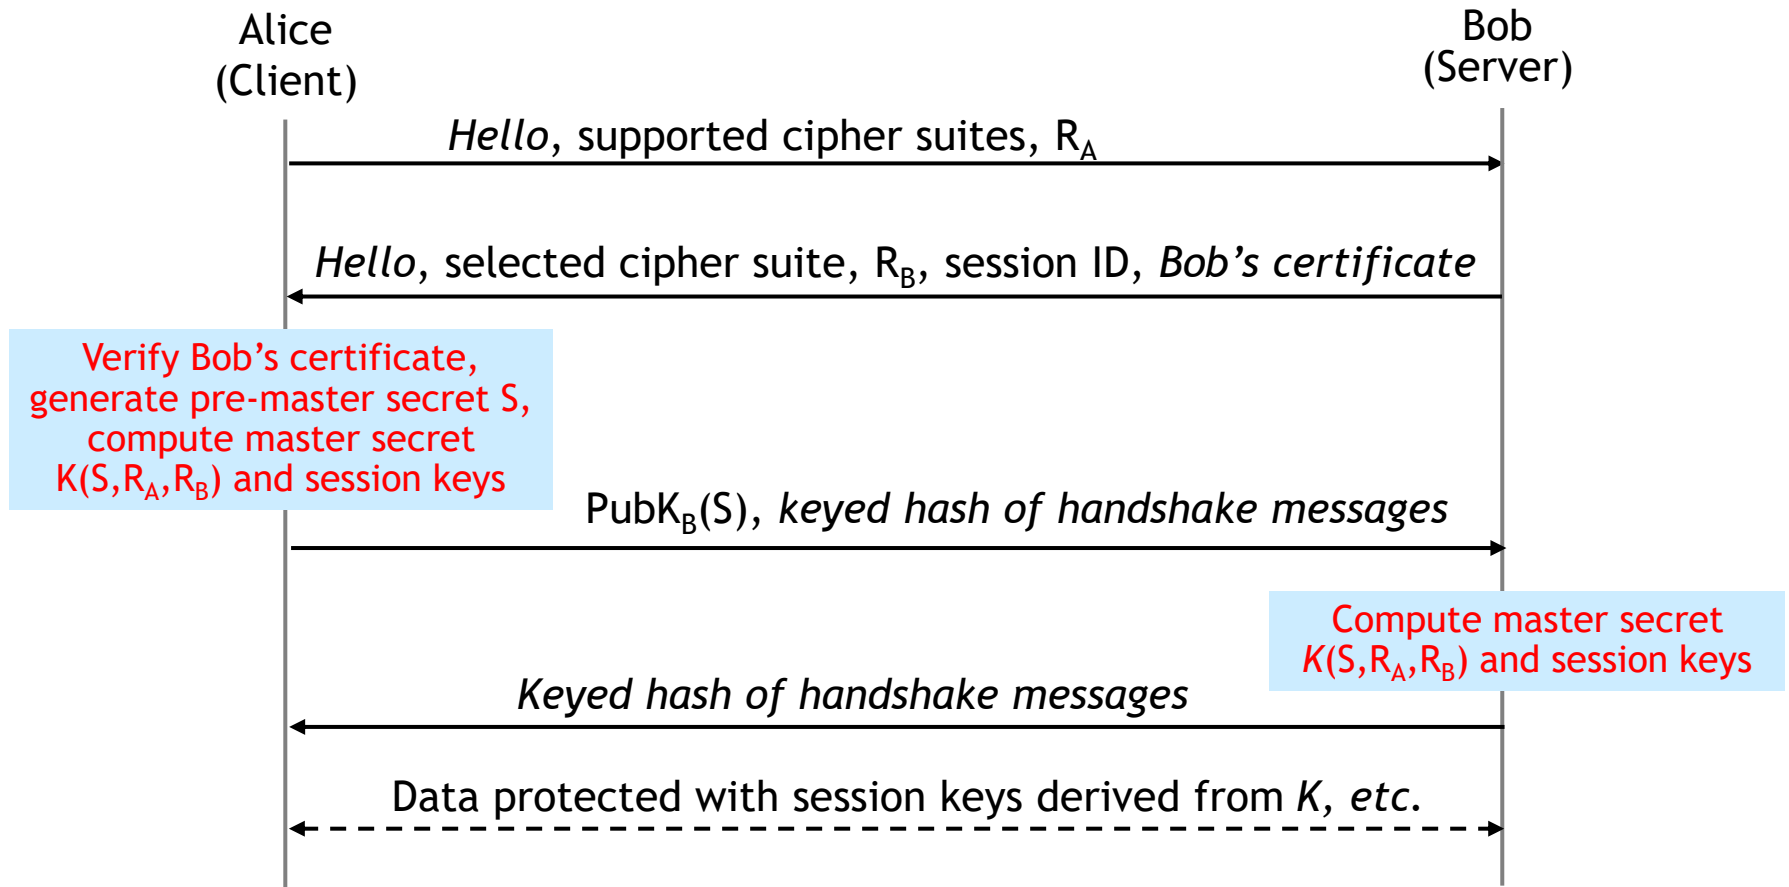
**Transport layer:** Reliable transport (e.g., TCP or SCTP)

# TLS RECORD PROTOCOL



Session keys and algorithms for compression, encryption, and MAC are negotiated during the TLS handshake

# OUTLINE OF A TYPICAL TLS SESSION SETUP



Only Bob is authenticated (based on the certificate).  
Alice needs to be authenticated at the application layer.

A cipher suite (covering key exchange, encryption and MAC) is negotiated.

# AUTHORIZATION PROTOCOLS



	<i>Object 1</i>	<i>Object 2</i>	<i>Object 3</i>	<i>Object 4</i>
<i>Subject 1</i>	read, write, execute	read, write, execute	read	read
<i>Subject 2</i>	read, execute	read, execute	null	null
<i>Subject 3</i>	read, execute	read	read, write	write

## ACCESS CONTROL MATRIX



# CAPABILITY LISTS

*Subject 1* → Object 1: read, write, execute; Object 2: read, write, execute; Object 3: read; Object 4: read

*Subject 2* → Object 1: read, execute; Object 2: read, execute

*Subject 3* → Object 1: read, execute; Object 2: read; Object 3: read, write; Object 4: write

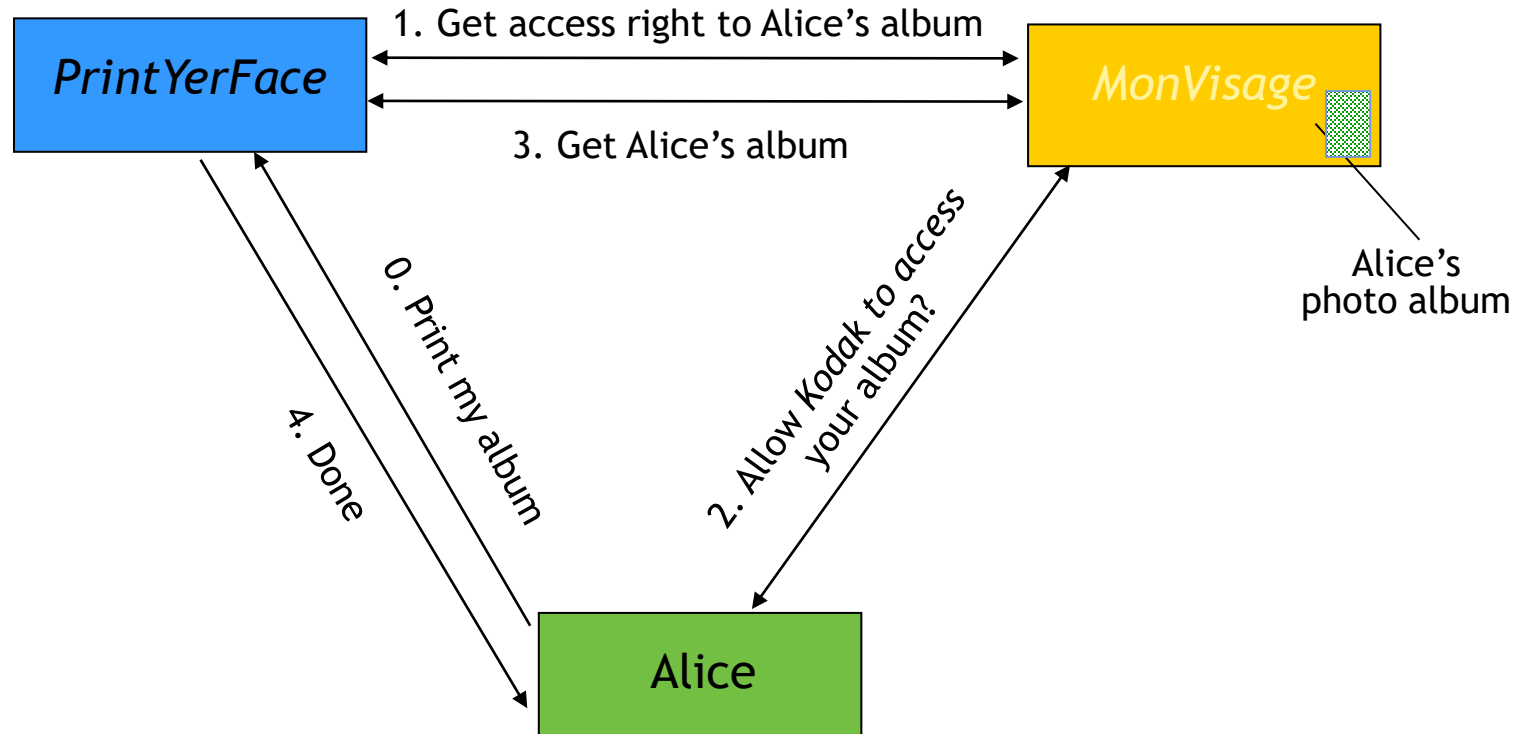


# *OAuth* DEVELOPMENT HISTORY

- The community effort (<http://oauth.net/>) was inspired by OpenID and started in late 2006
- The IETF formed a working group (<http://datatracker.ietf.org/wg/oauth/>) to standardize *OAuth* in May 2009
  - It has published RFC 5849 (*Informational*) based on *OAuth* Core 1.0 Revision A
  - Current focus is on development of *OAuth* 2.0 taking into account RFC 5849 and the Web Resource Authorization Protocol (WRAP)
  - The effort has an expanding support base, including telecom operators and over-the-top service providers

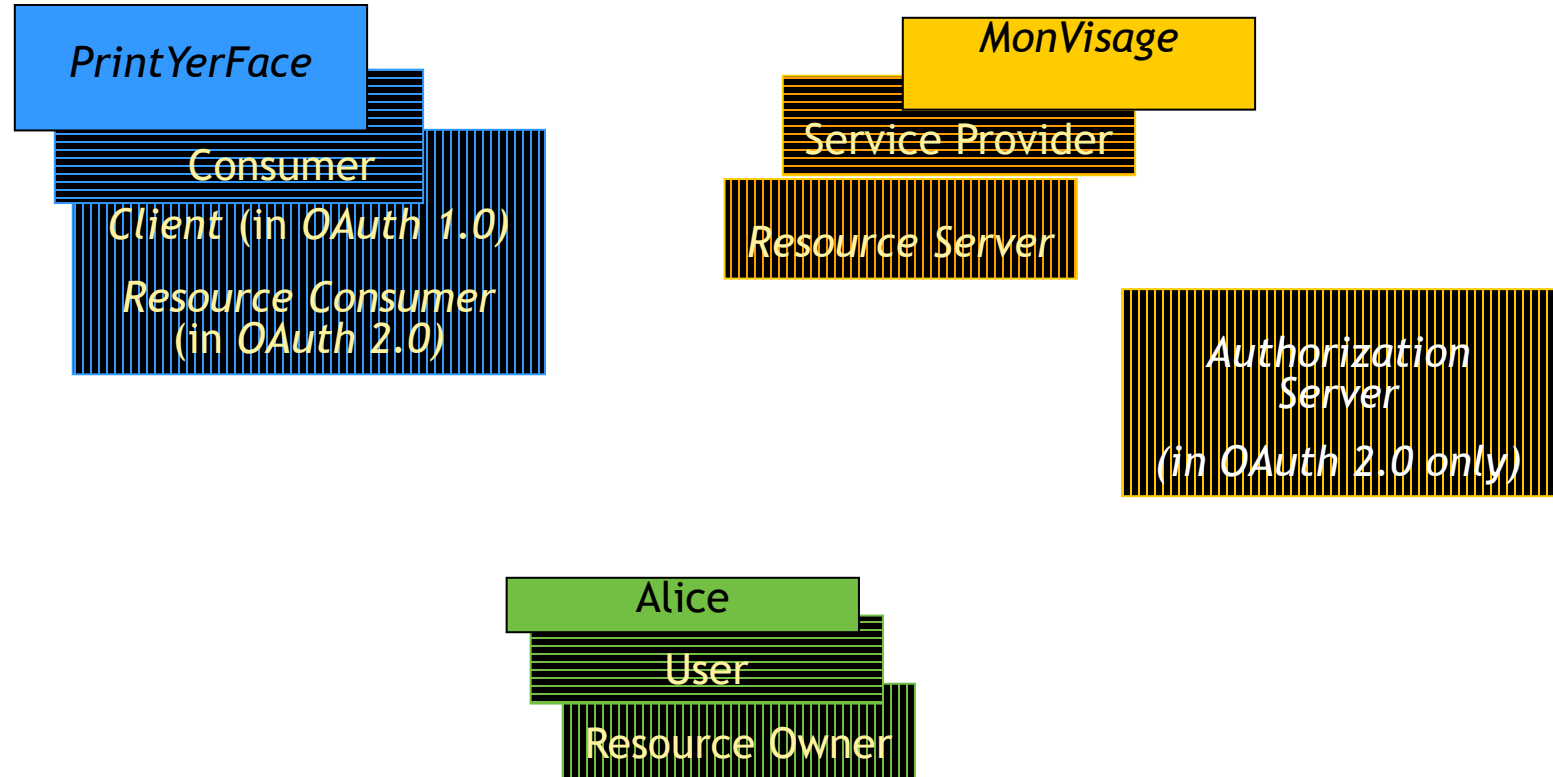


# OPEN AUTHORIZATION (*OAUTH*) PROTOCOL



Alice grants *PrintYerFace* access to her album at *MonVisage* without revealing her credentials

# *OAUTH* TERMINOLOGY

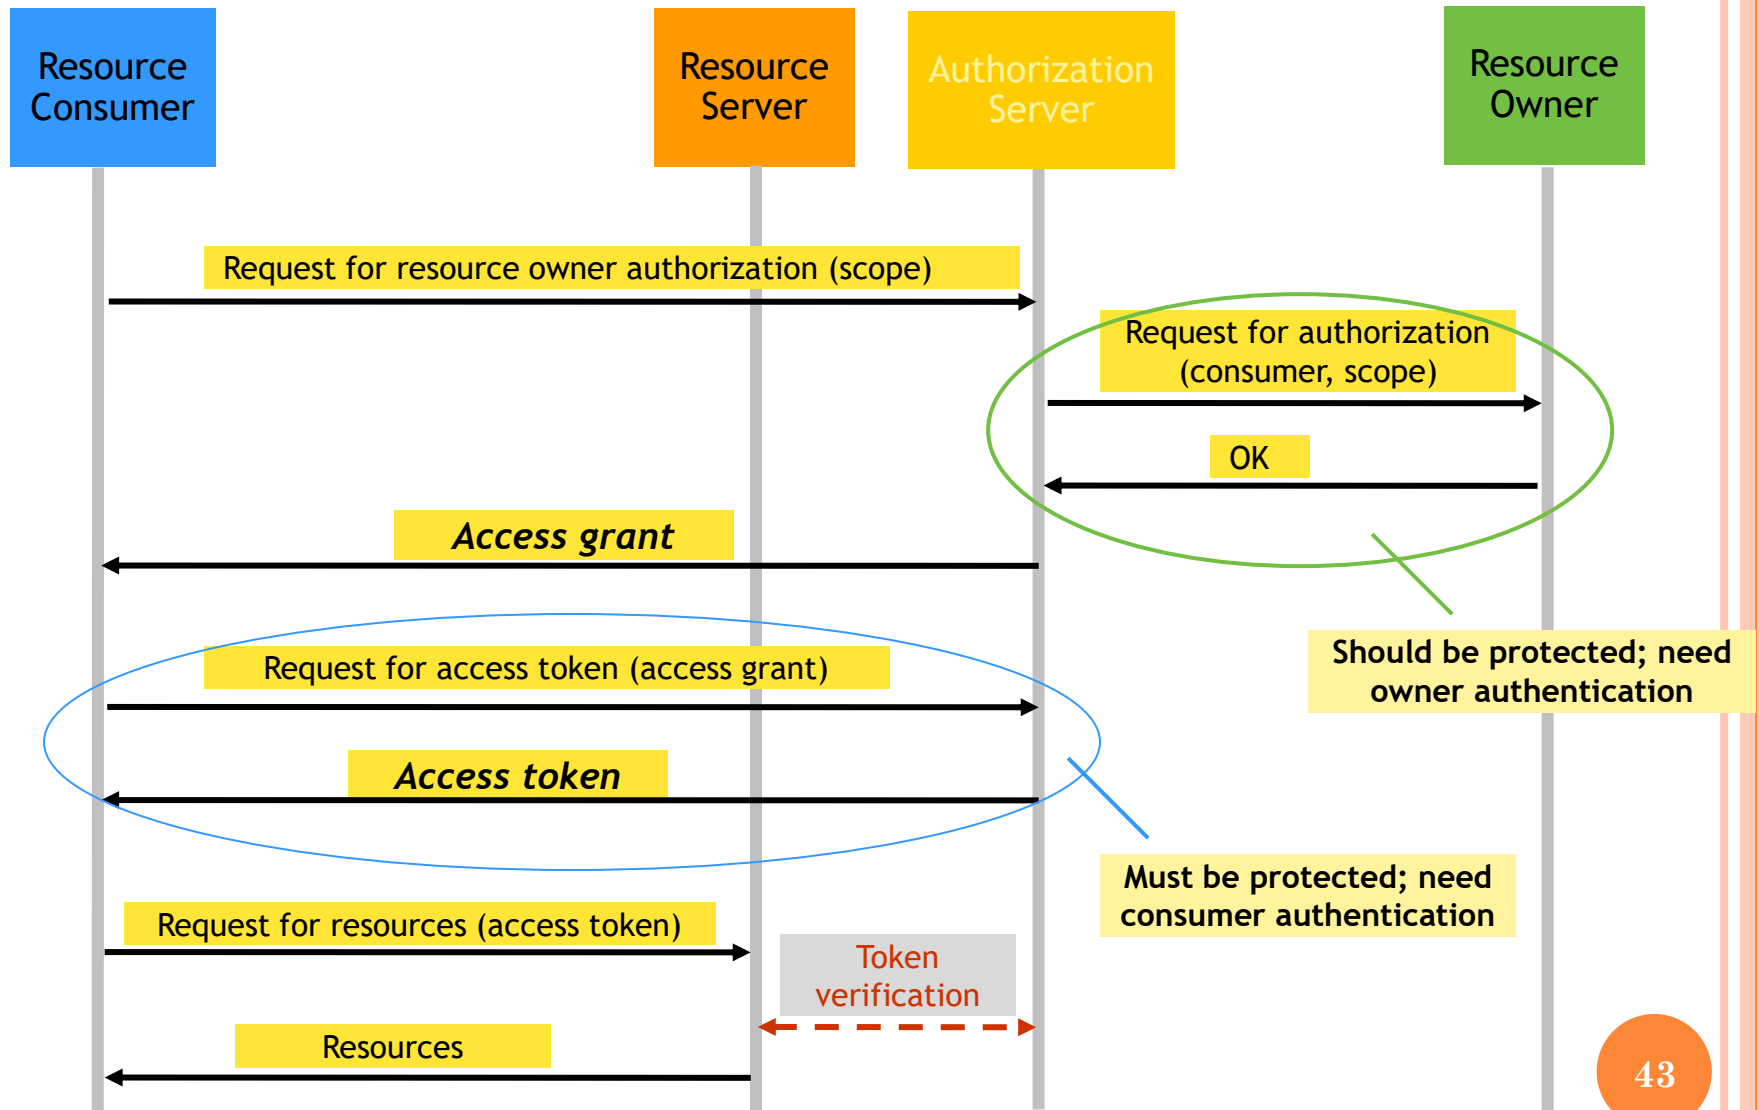


community term



IETF term

# *OAUTH* 2.0 CONCEPTUAL PROTOCOL VIEW



# IDENTITY MANAGEMENT IN THE CLOUD



# PROTECTION OF *DATA AT REST*

- In SaaS and PaaS, protection for data is expected to be performed by the Cloud *provider*
- In IaaS, protection of data is also an essential responsibility of the Cloud *consumer*
- Issues to consider for the decision on where to maintain encryption keys:
  - access control (authorization)
  - availability of the keys,
  - compliance (whether it is necessary to ***escrow*** the keys with the provider)

**Case study: Amazon's S3 storage supports encryption options and provides an option for key escrowing**

# PROTECTION OF *DATA IN MOTION*

- Access to data
  - End-user authentication (including key establishment)
  - End-user authorization for accessing specific records
- Data transfer among VMs
- Data transfer among data centers
  - Secure tunnels (possibly provided by the *Network as a Service* provider)

# CLOUD DATA STORAGE CONSIDERATIONS

- Can an enterprise store *all* its data in the Cloud?
  - Who (the enterprise or Cloud provider) is responsible for the private or business-sensitive information?
  - How to ensure compliance with the regulations?
  - Does the Cloud Provider guarantee data security “*at rest*” and “*in motion*”?
  - Who (enterprise or Cloud provider) is responsible for key distribution and management?
  - Who is responsible for scheduling and carrying back-ups?
- What happens if the Cloud provider goes out of business?

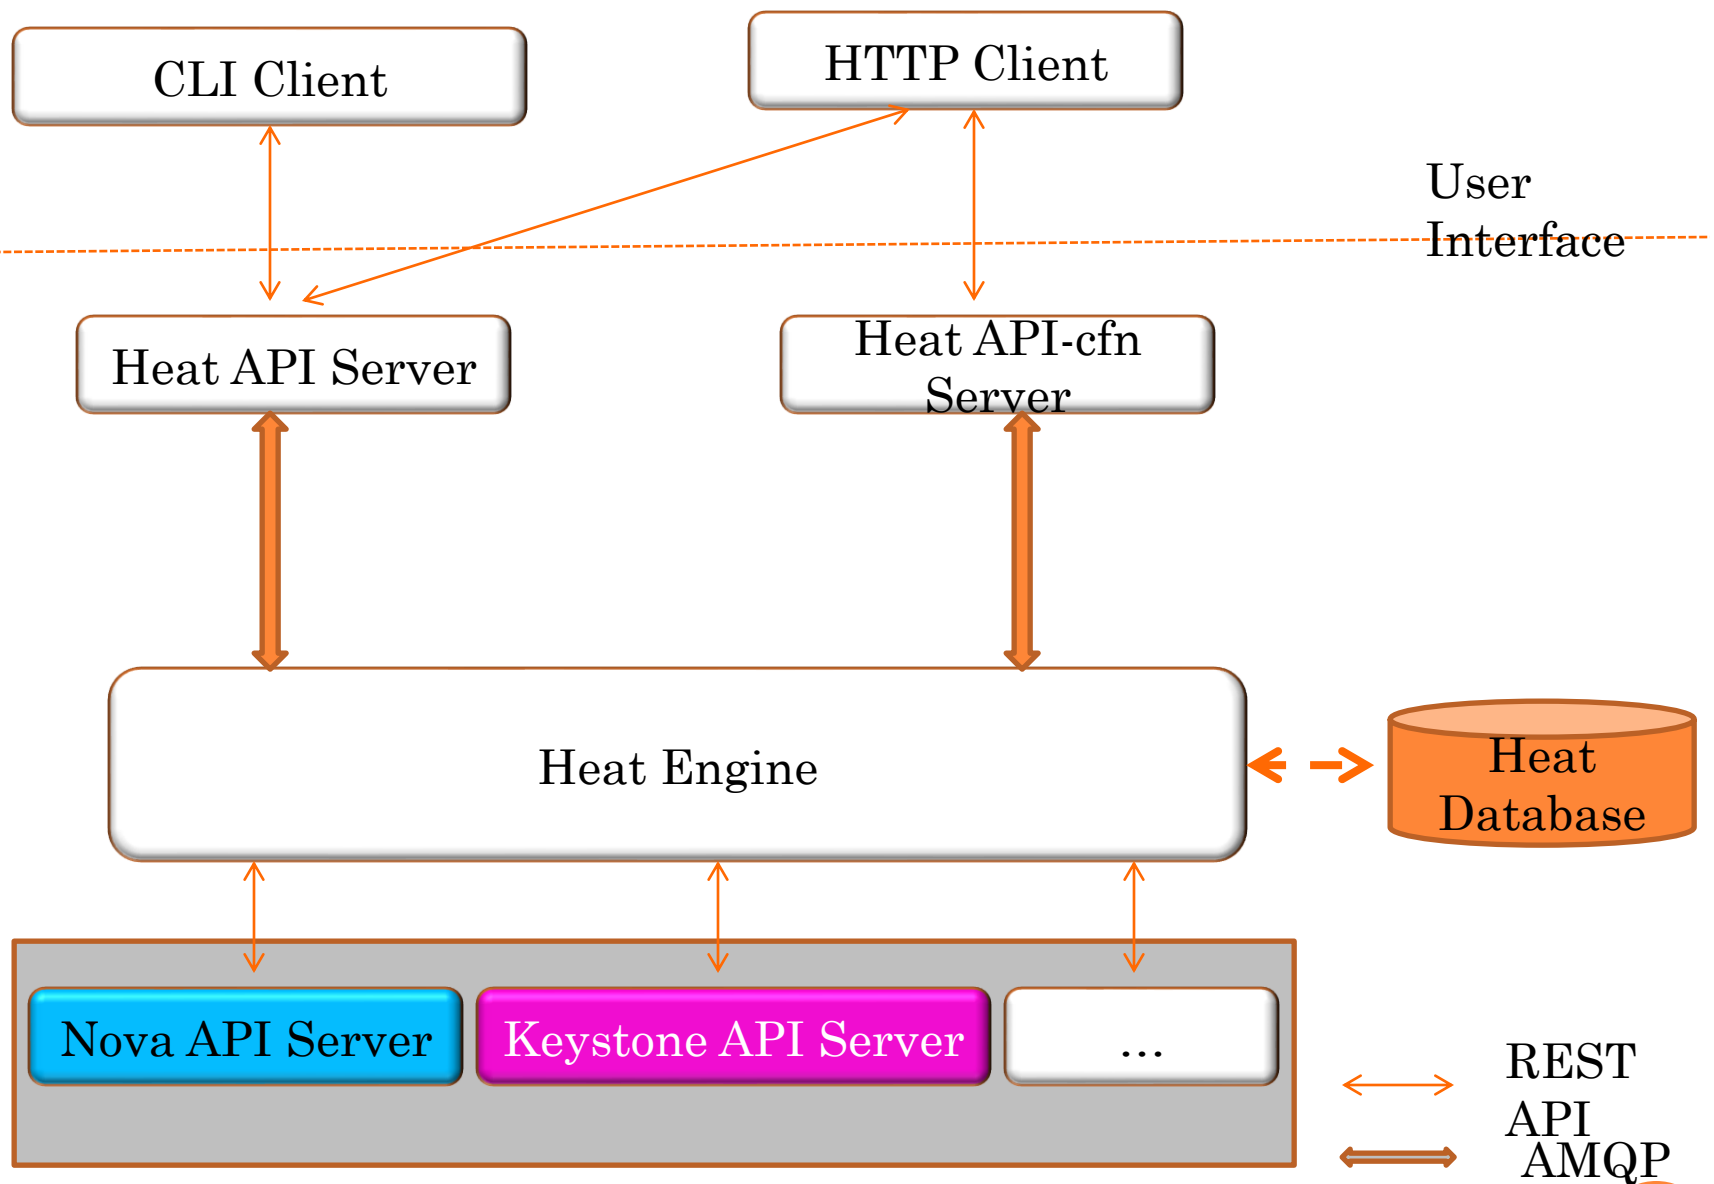
The data storage problem should never have a single *all-or-nothing* solution!

# IDENTITY MANAGEMENT IN OPENSTACK

- *Keystone* is the *gatekeeper*
- *Keystone*
  - provides centralized authentication and authorization service
  - supports the constructs of domain, project, group, and role
  - generates and validates *tokens*

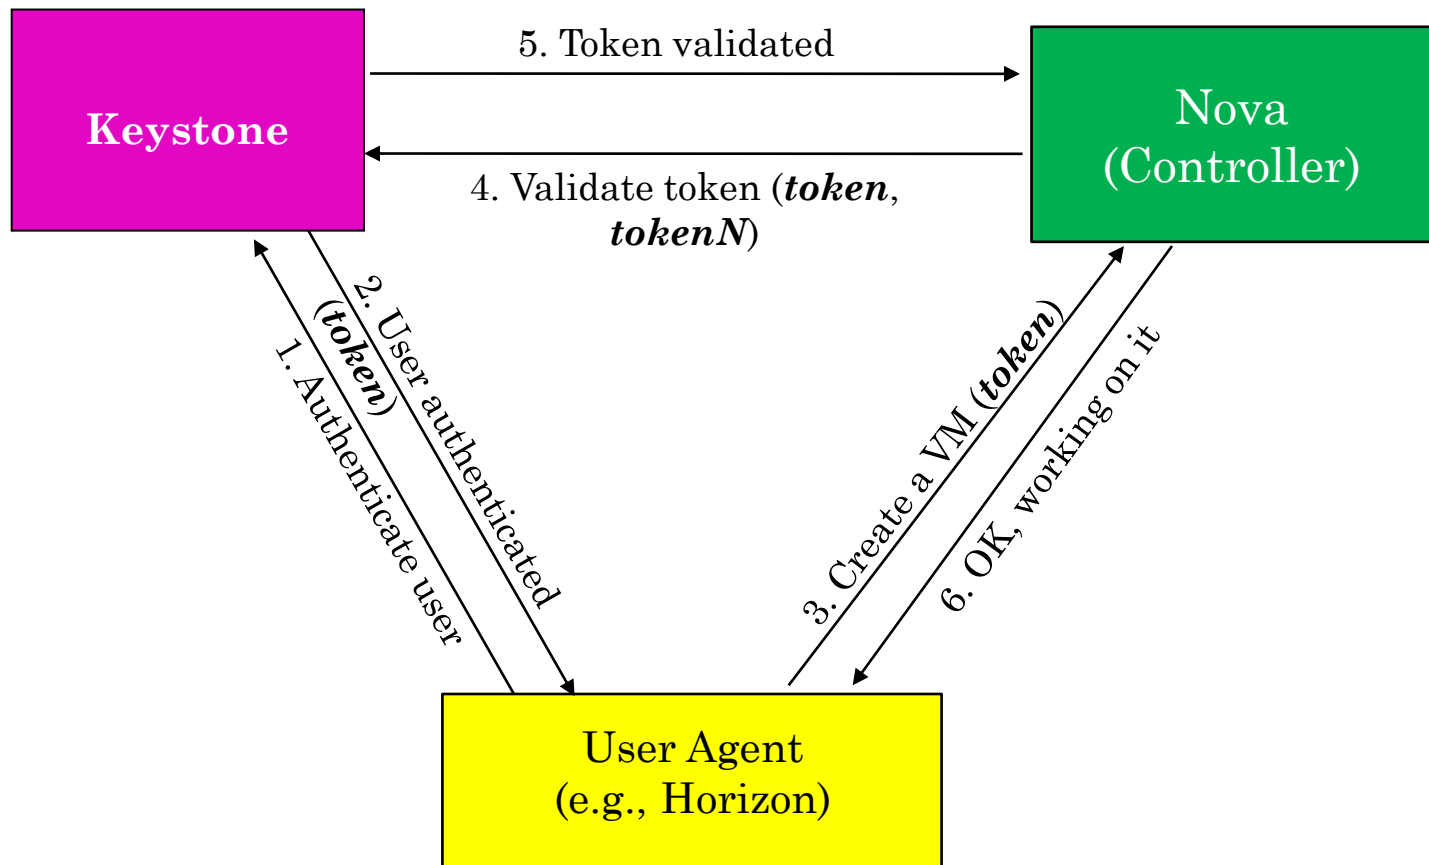




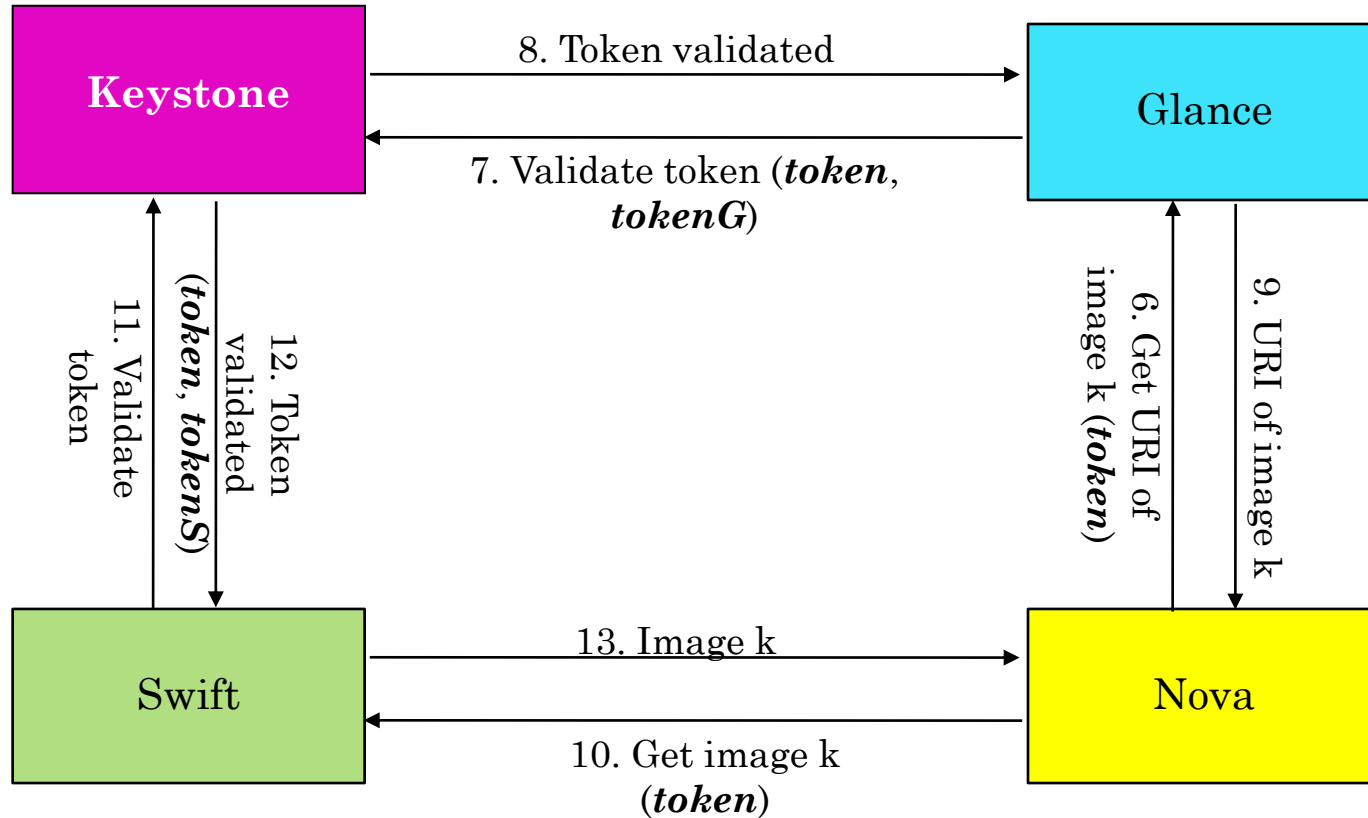


The *Heat* computing architecture

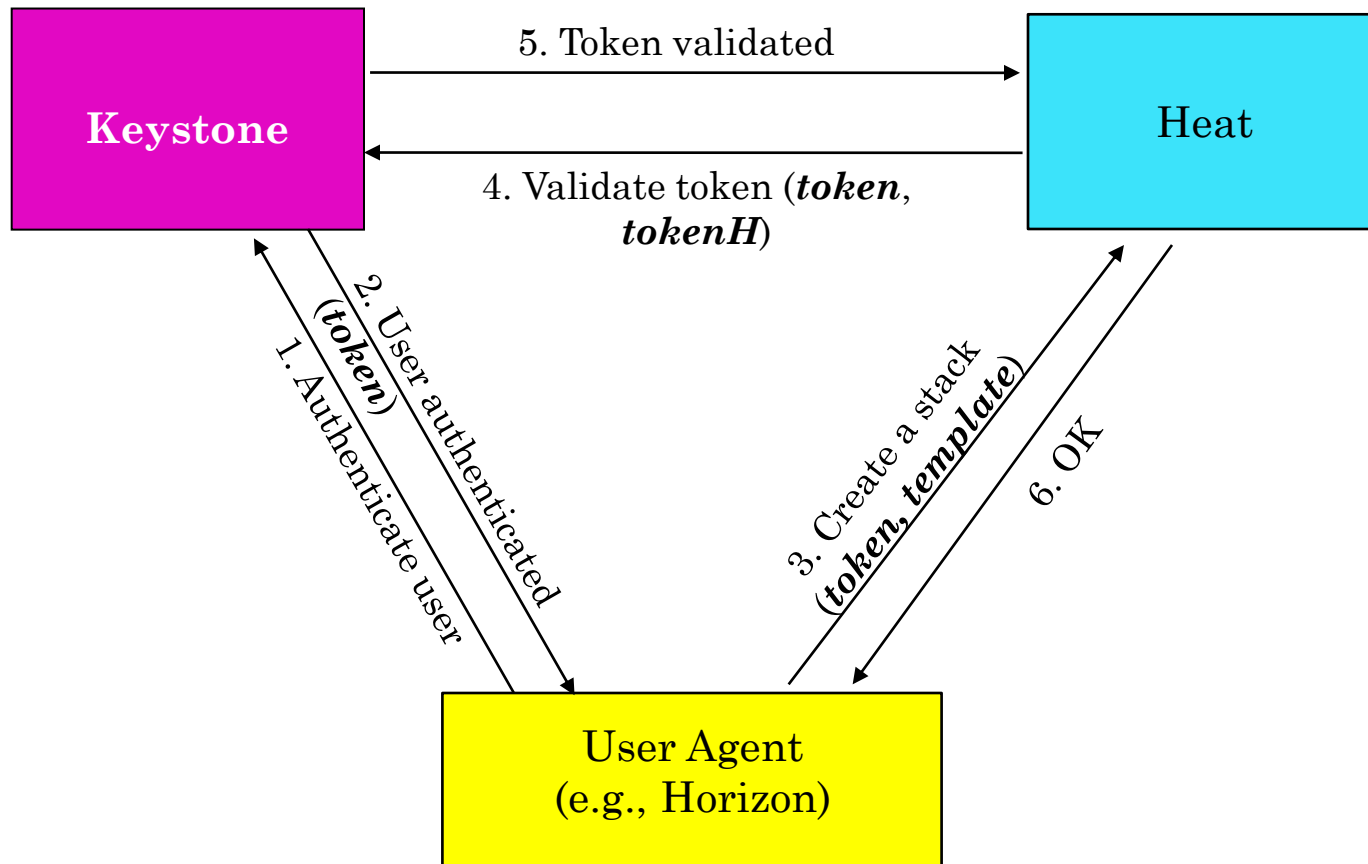
# A SIMPLIFIED WORKFLOW FOR VM PROVISIONING IN OPENSTACK



## ADDITIONAL STEPS FOR VM PROVISIONING



## A SIMPLIFIED WORKFLOW FOR AUTO-SCALING

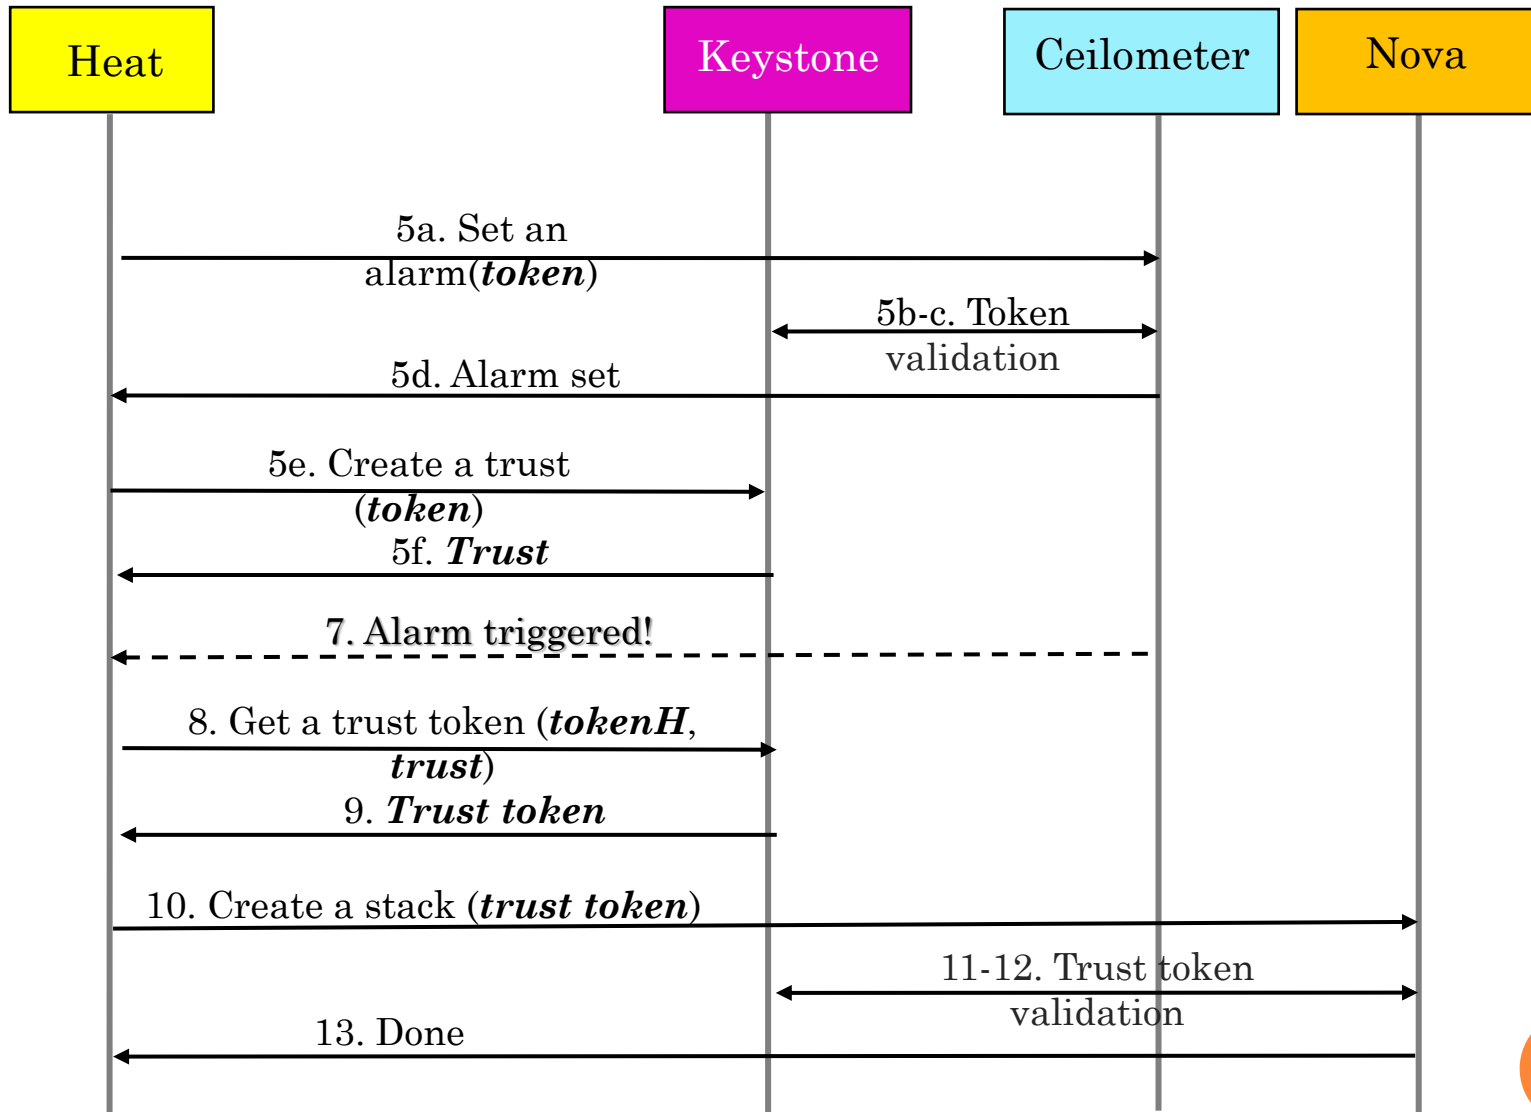


## AN EXAMPLE TOKEN (UNSIGNED)

```
{
  "token": {
    "expires_at": "2015-05-27T22:52:58.852167Z",
    "issued_at": "2015-05-27T21:52:58.852167Z",
    "methods": ["password"],
    "domain": {
      "id": "3a5140aecd974bf08041328b53a62458",
      "name": "Wonderland"
    },
    "roles": [{
      "id": "9fe2ff9ee4384b1894a90878d3e92bab",
      "name": "admin"
    }],
    "user": {
      "domain": {
        "id": "3a5140aecd974bf08041328b53a62458",
        "name": "Wonderland"
      },
      "id": "3ec3164f750146be97f21559ee4d9c51",
      "name": "Alice"
    }
  }
}
```



## ADDITIONAL STEPS FOR AUTO-SCALING

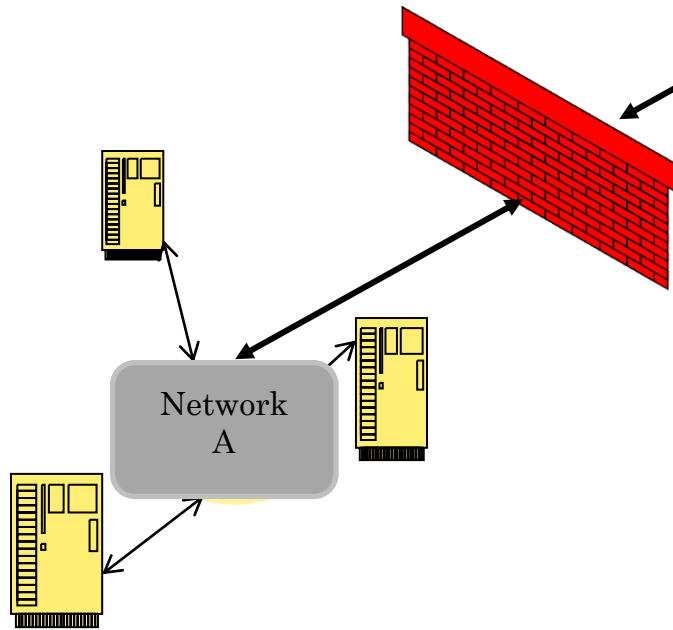


# AN EXAMPLE TRUST

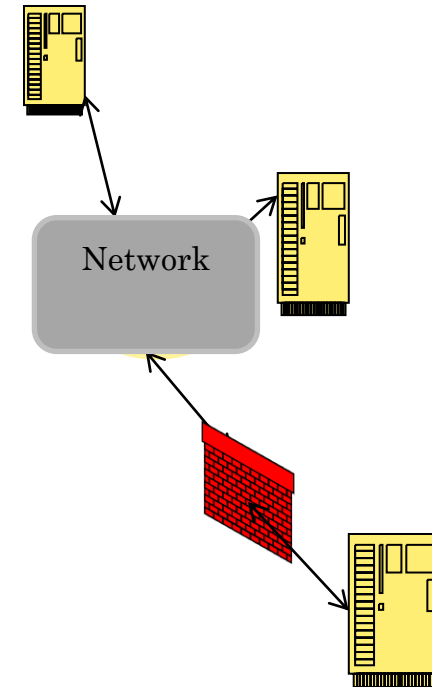
```
{
  "trust": {
    "expires_at": "2016-05-27T21:52:58.852167Z",
    "id": "c703057be878458588961ce9a0ce686b",
    "impersonation": true,
    "project_id": "fb49a0ecd60c4d2092643b4cfe272106",
    "remaining_uses": null,
    "roles": [{
      "id": "9fe2f79ee4384b1894a9083bd3e92bab",
      "name": "admin"}
    ],
    "trustee_user_id": "29beb2f1567642eb810b042b6719ea88",
    "trustor_user_id": "3ec3164f750146be97f21559ee4d9c51"
  }
}
```



# Firewalls



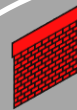
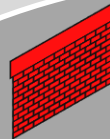
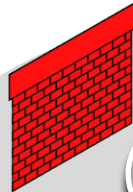
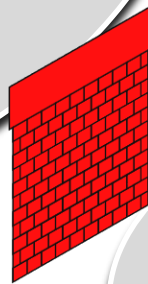
a) A firewall between two networks



b) A firewall protecting a single host



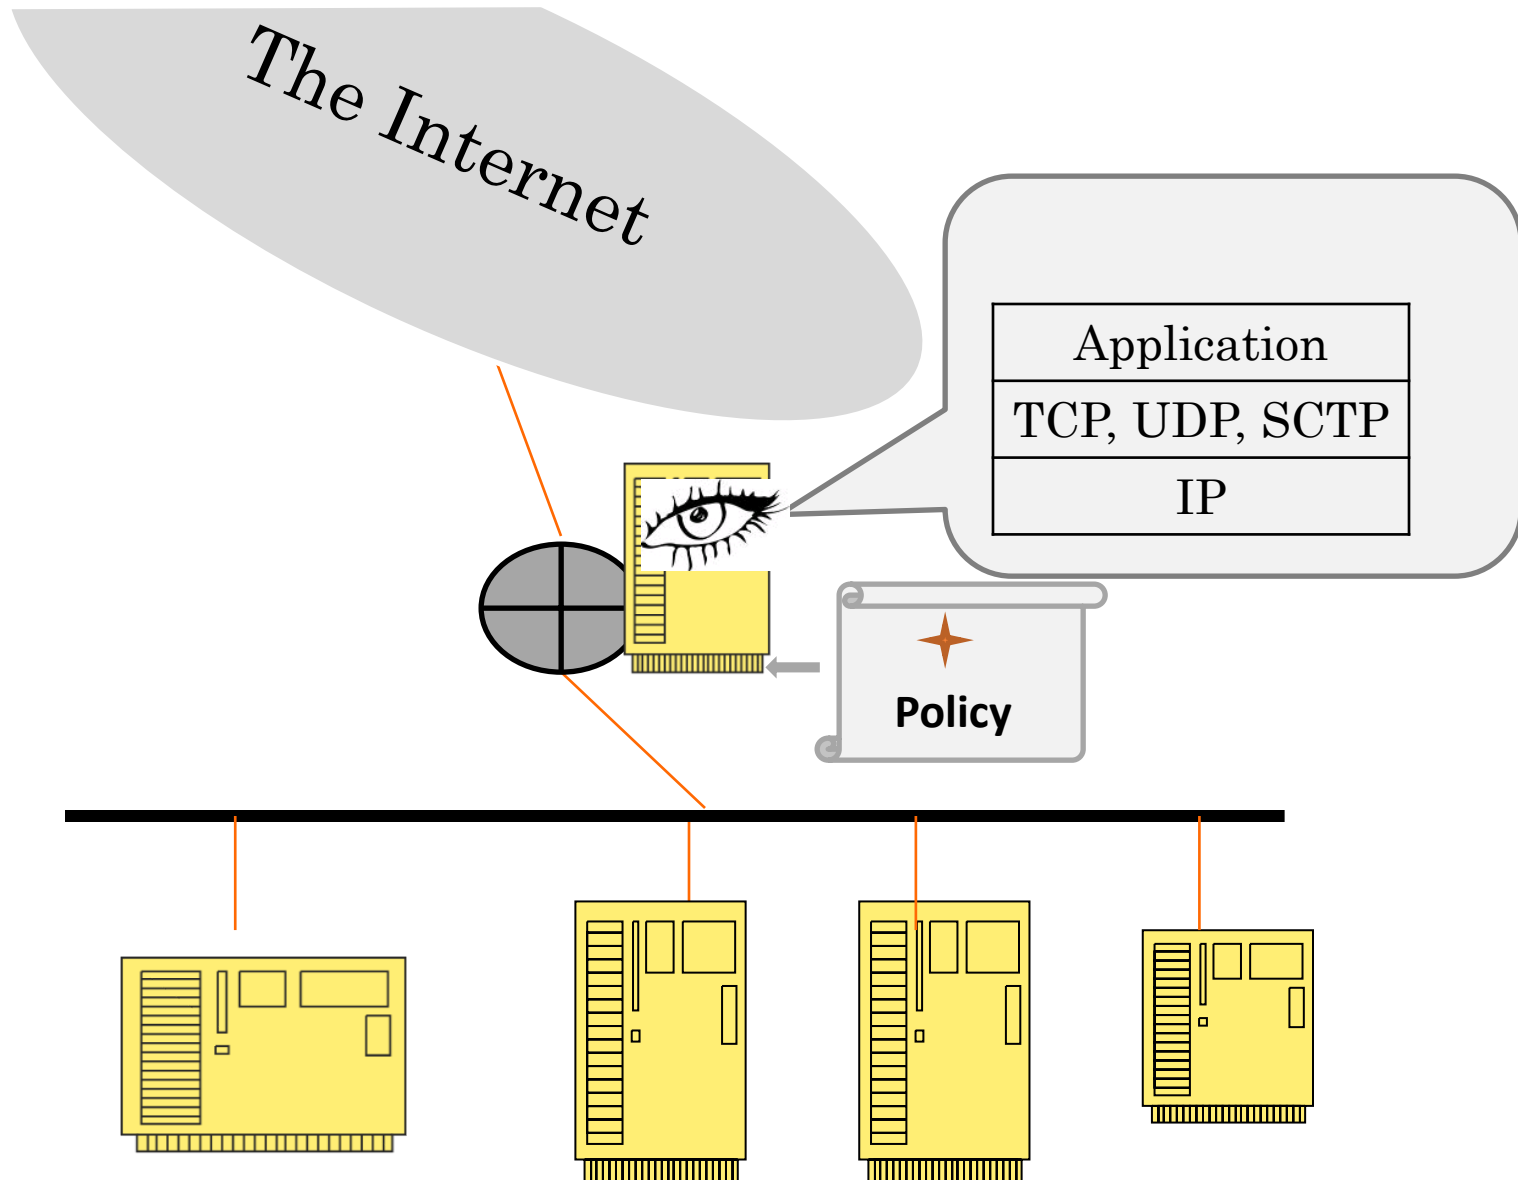
Public Network



Private Network

**Interconnecting networks with different security postures**

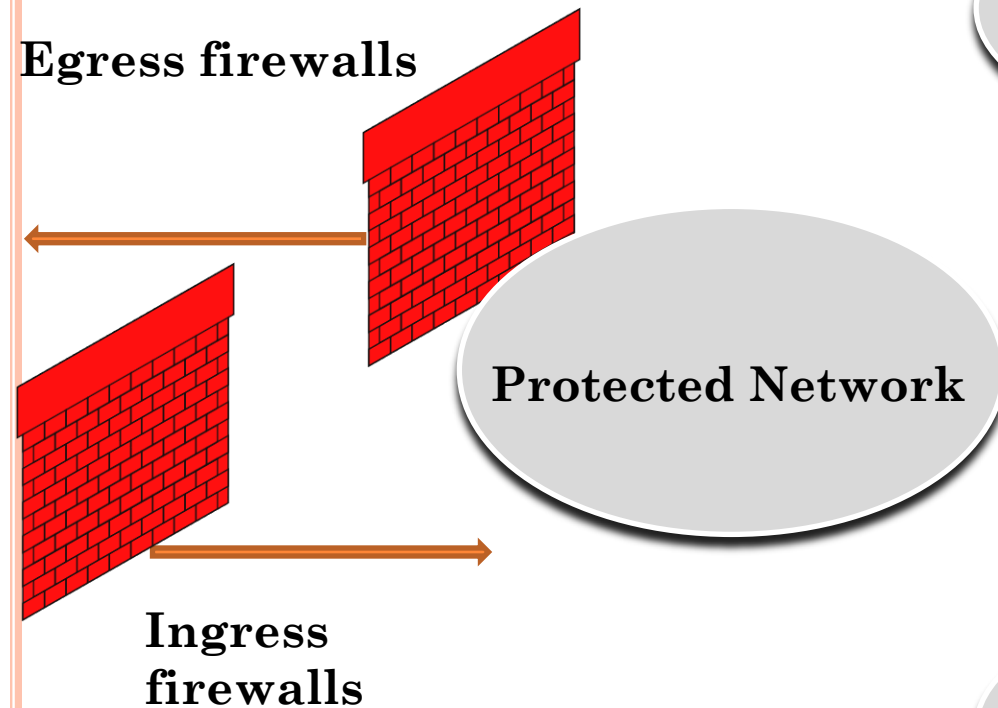




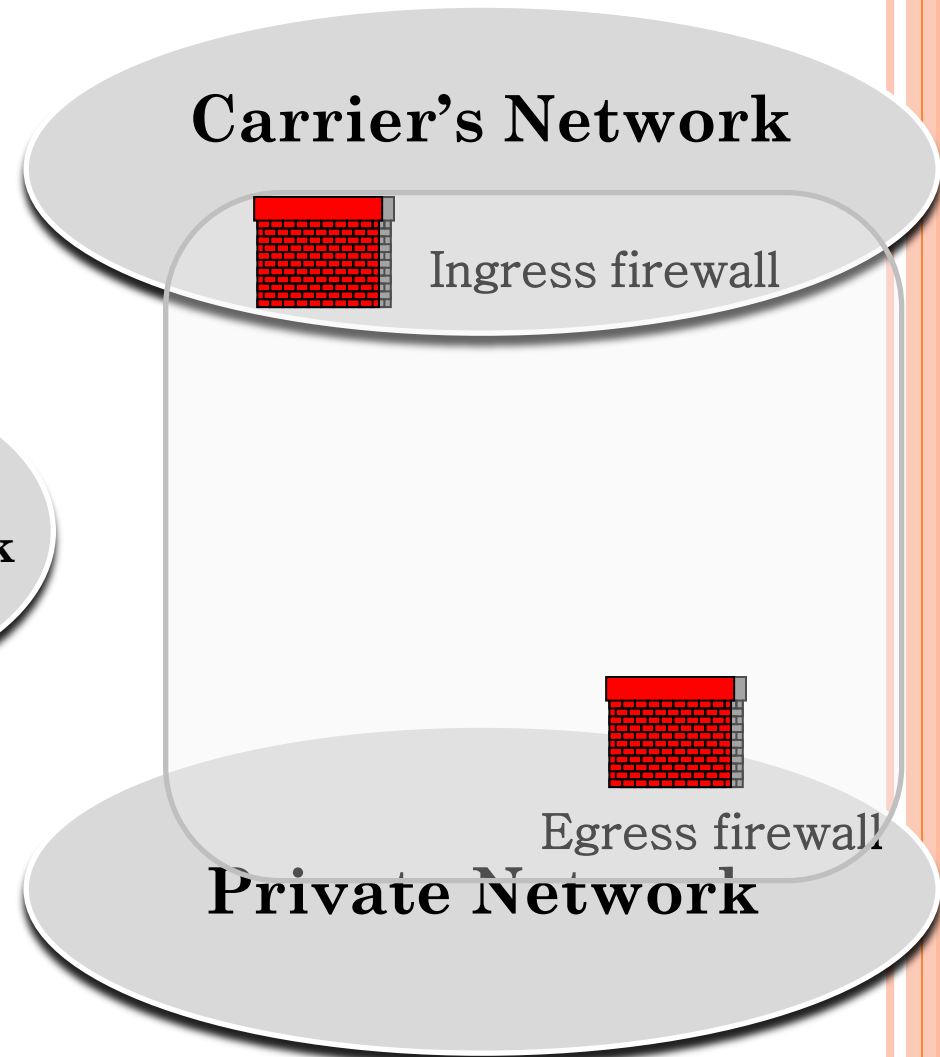
**The application gateway**



# Ingress and egress firewalls

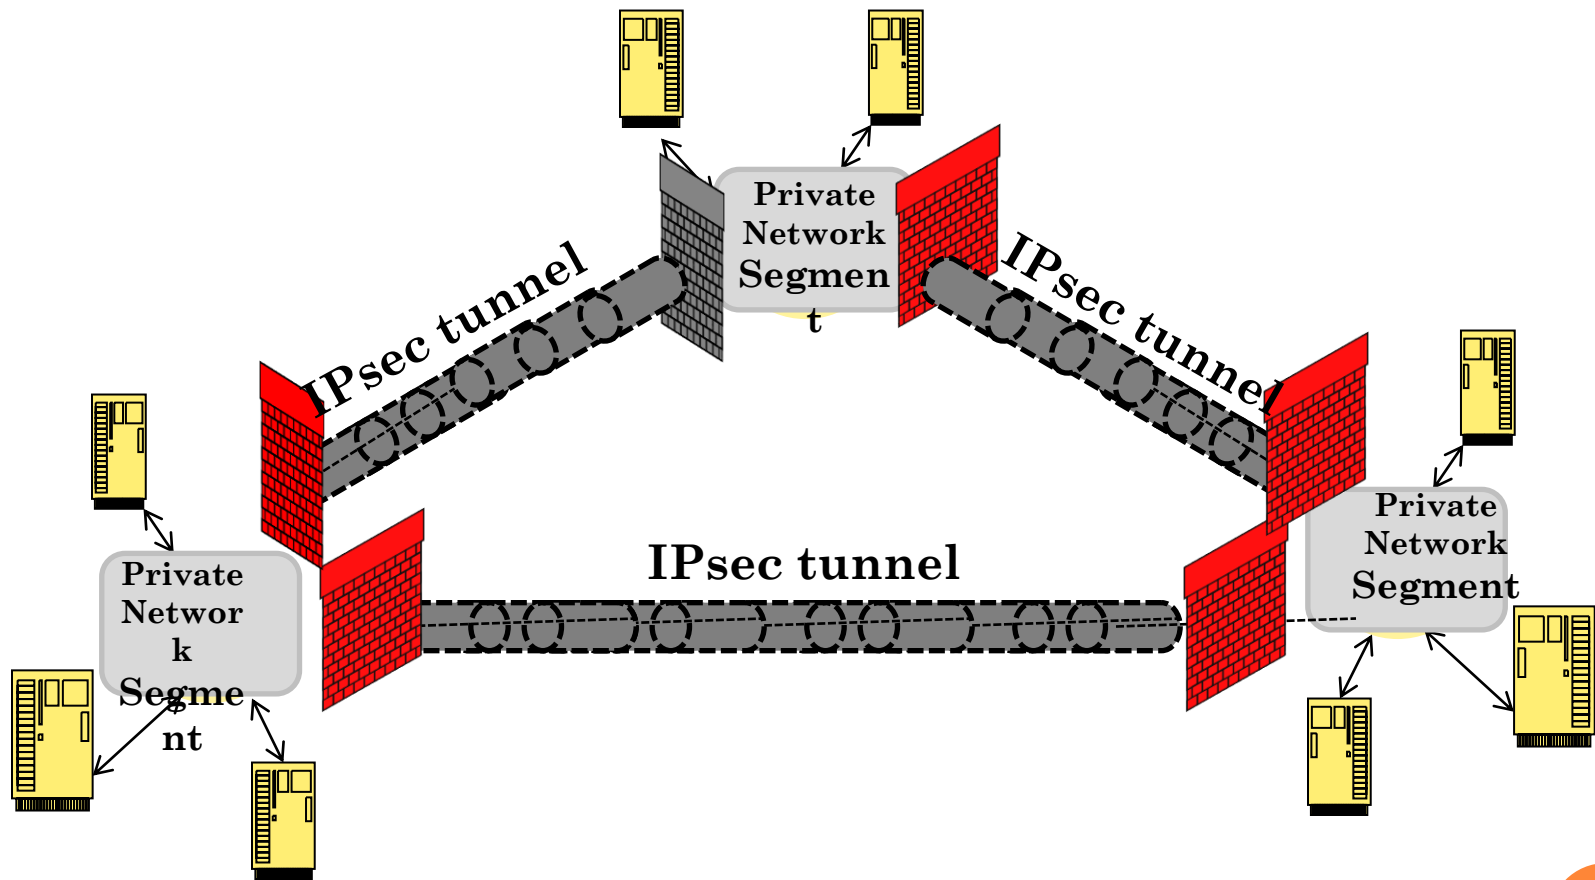


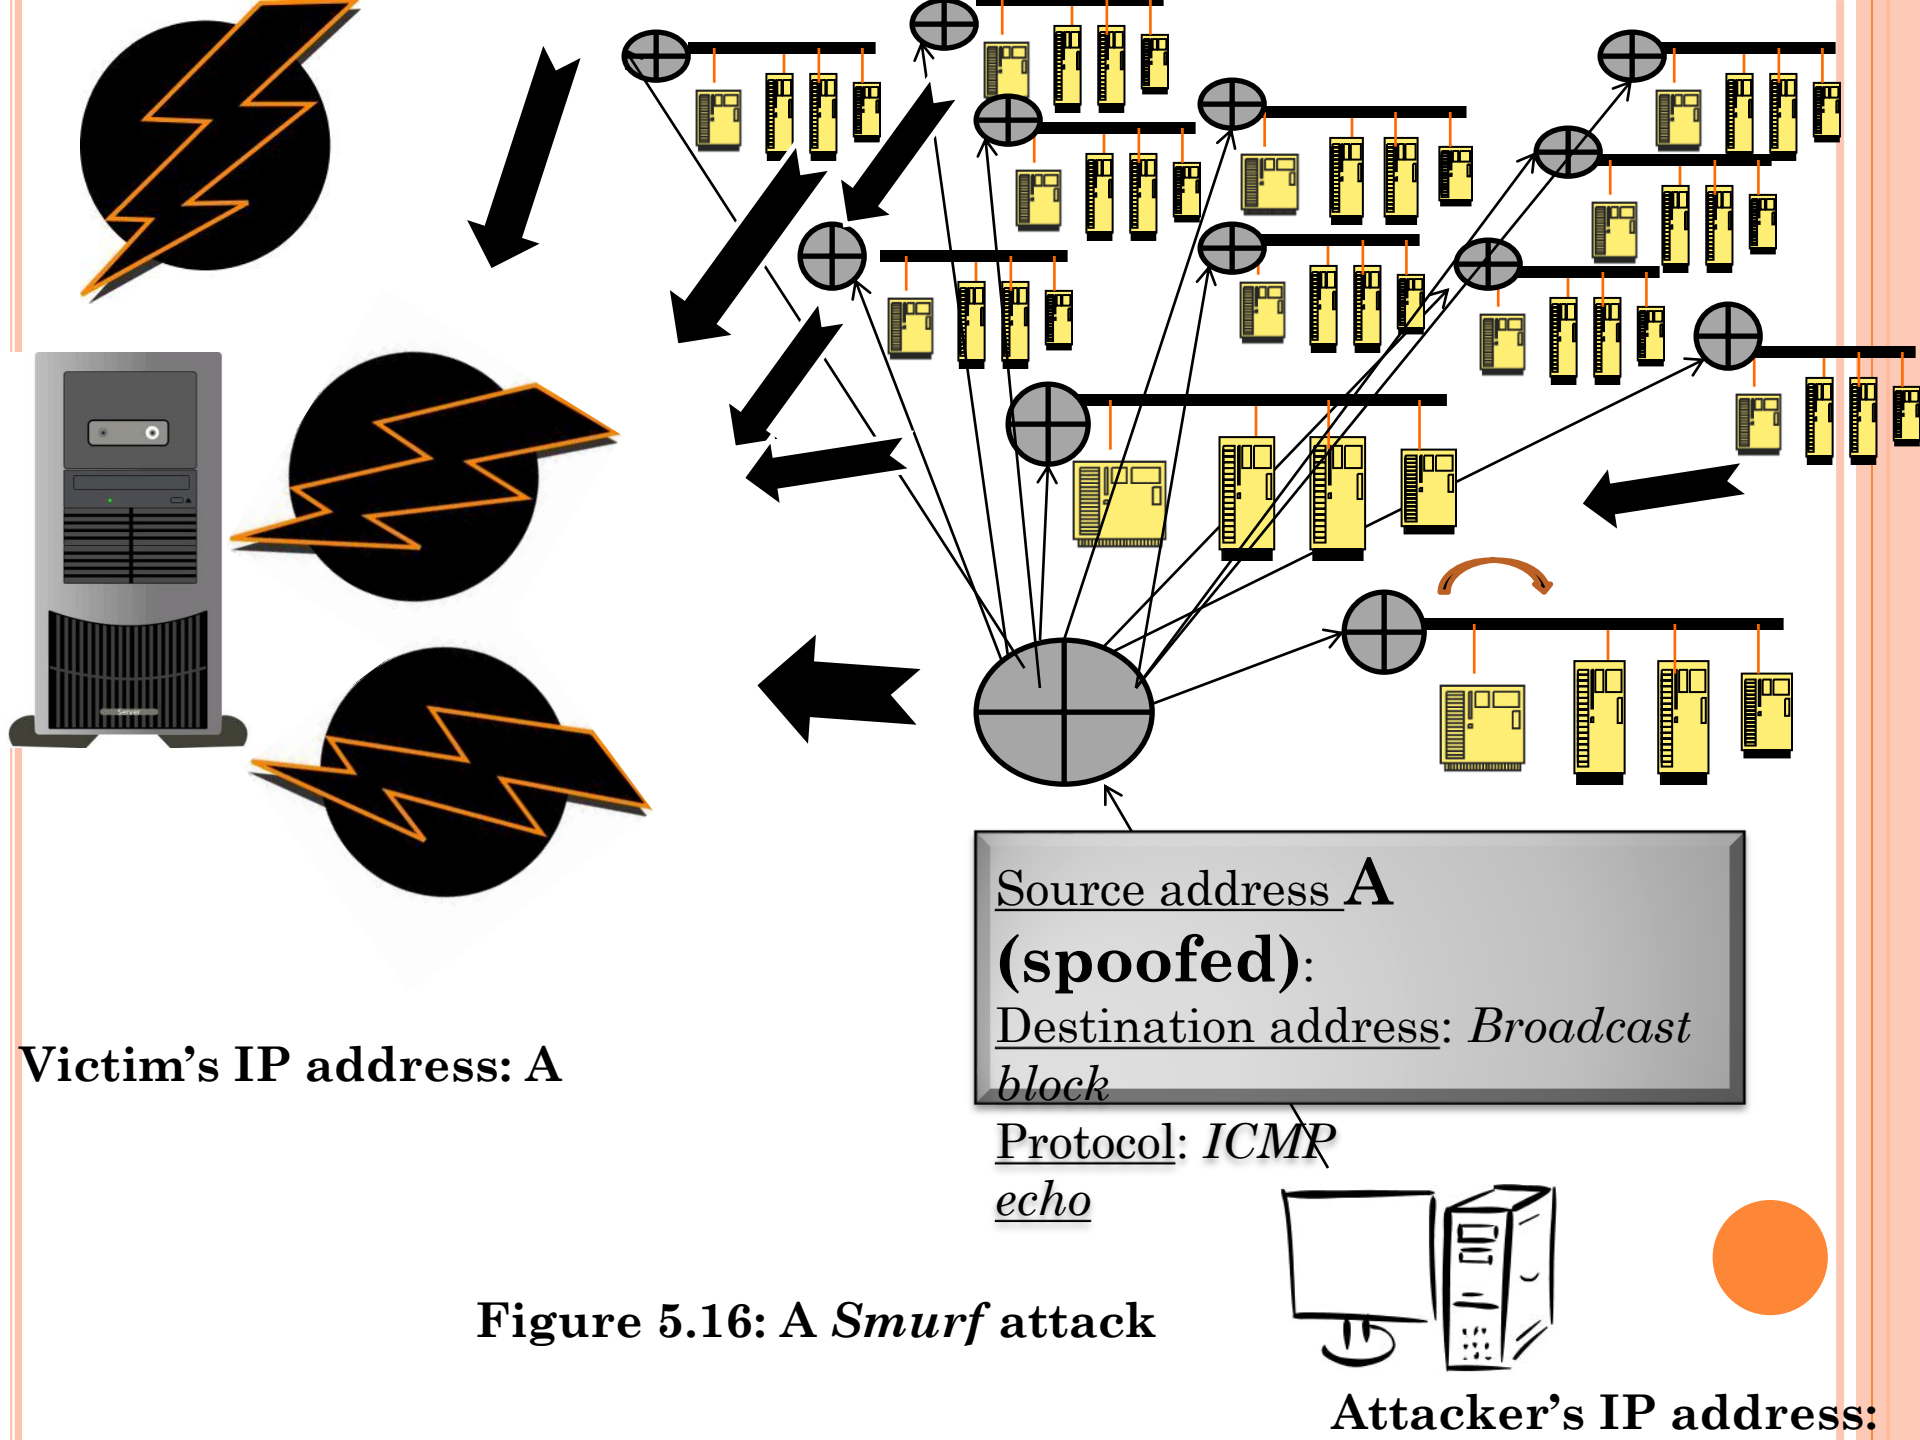
a) Interfaces



b) Split CPE

# Layer 3 VPN with firewalls





Victim's IP address: A

Source address A  
(spoofed):  
Destination address: *Broadcast block*  
Protocol: *ICMP echo*

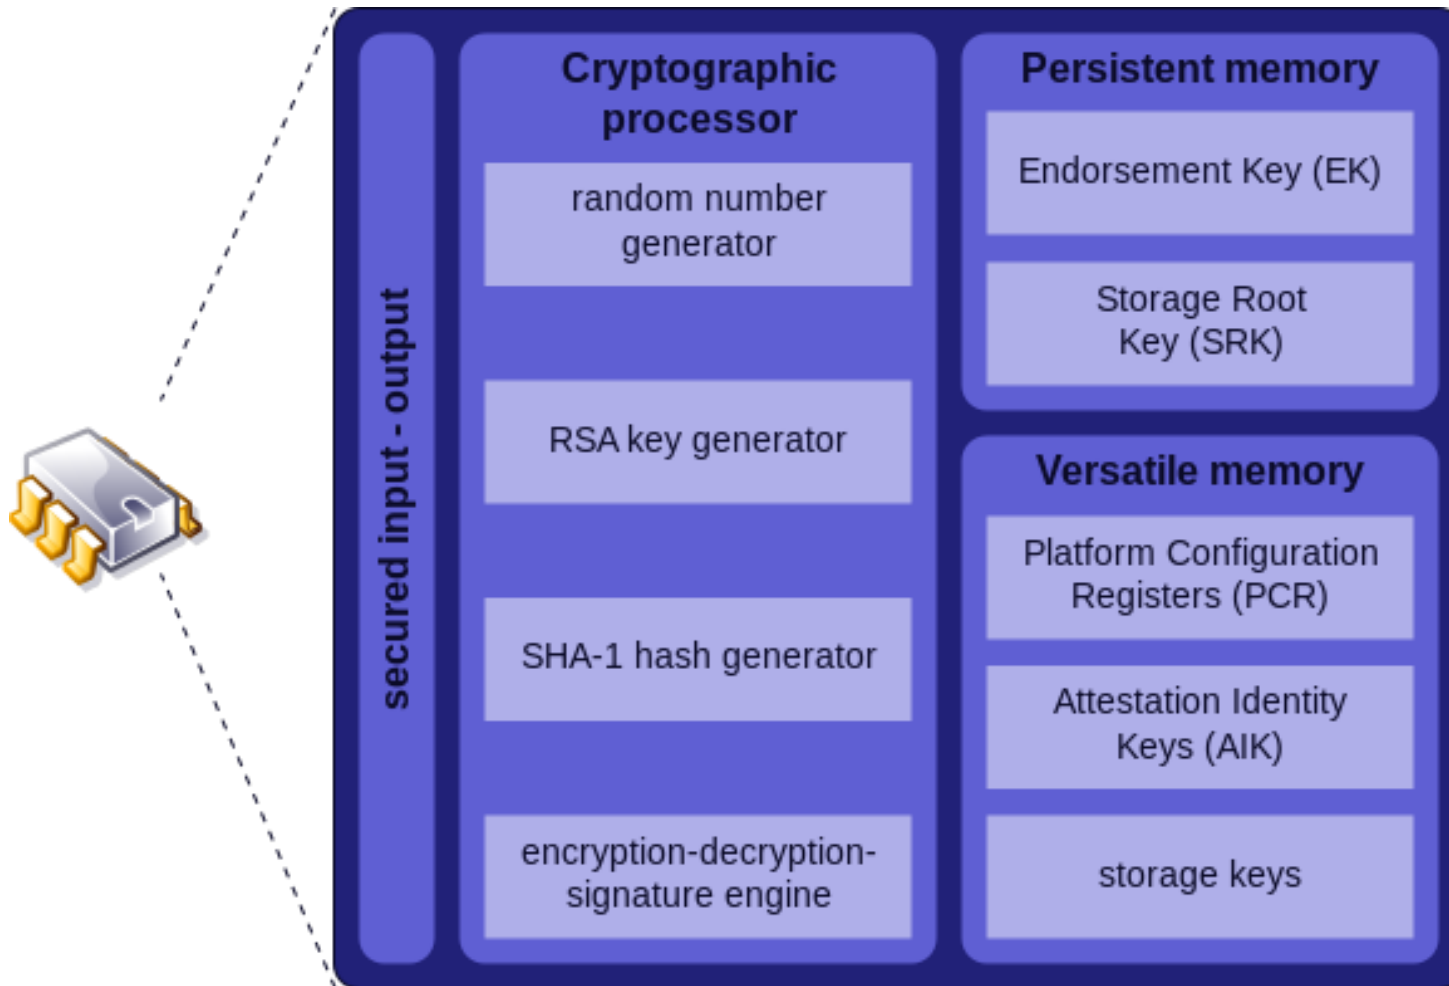
Figure 5.16: A *Smurf* attack



Attacker's IP address:

# TRUSTED PLATFORM MODULE (TPM)

- A standard for a secure cryptoprocessor
  - Specified by the Trusted Computing Group (TCG).
  - Standardized as ISO/IEC 11889
- Major capabilities:
  - **Remote attestation** – via an unforgeable hash key summary of the hardware and software configuration to allow a third party to verify that the software has not been changed.
  - **Binding** – encrypts data using TPM bind key, a unique RSA key descended from a storage key
  - **Sealing** – binding, to which a required state of the TPM to be at so as to unseal the data is specified



Source: Eusebius (*Guillaume Piolle* ,  
<https://commons.wikimedia.org/w/index.php?curid=4809738>

# KEY APPLICATIONS OF TPM IN THE CLOUD

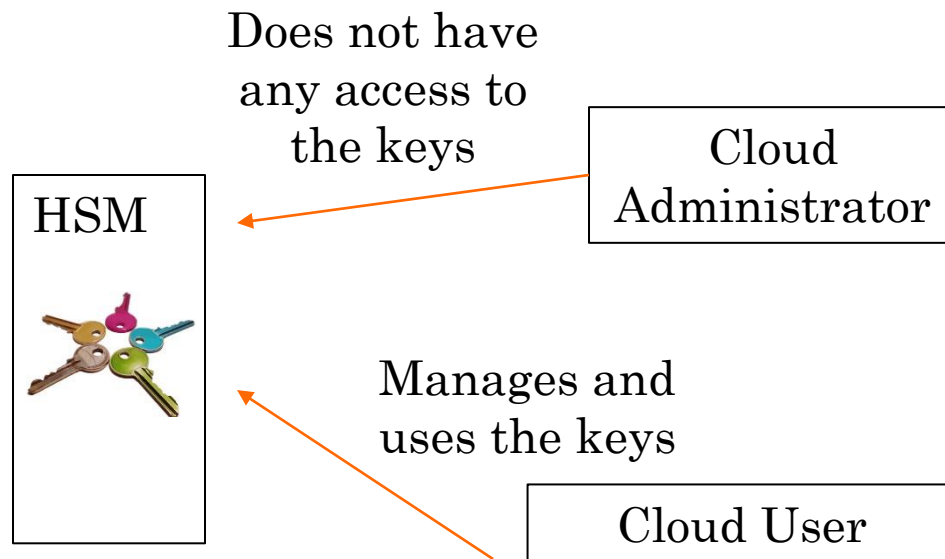
- **Ensuring platform integrity:** TPM forms a *root of trust*. The BIOS, hypervisors, and the OS build on this. The Intel Trusted Execution Technology (TXT) creates a *chain of trust* to attest remotely that a computer has a specified hardware setup and is using specified software..
- **Disk encryption**
- **Password protection**
- **Digital rights management**
- **Enforcement of software licenses**
- **Storing the root certificates**





# HARDWARE SECURITY MODULES

- A physical computing device that stores and manages cryptographic keys and also performs cryptoprocessing. Thus no key escrowing is needed!
- Unlike TPM, HSMs are plug-in devices—often pooled externally.



# THE LAST HOMEWORK

- Read and understand <https://aws.amazon.com/cloudhsm/details/>
- Close all Amazon instances that you have created lest you be charged!

